

On The Landscape of Spoken Language Models: A Comprehensive Survey

Siddhant Arora^{1*} Kai-Wei Chang^{2*} Chung-Ming Chien^{3*} Yifan Peng^{1*} Haibin Wu^{2*#}
Yossi Adi⁴⁺ Emmanuel Dupoux⁵⁺ Hung-Yi Lee²⁺ Karen Livescu³⁺ Shinji Watanabe¹⁺

¹ Carnegie Mellon University, USA

² National Taiwan University, Taiwan

³ Toyota Technological Institute at Chicago, USA

⁴ Hebrew University of Jerusalem, Israel

⁵ ENS - PSL, EHESS, CNRS, France

Abstract

The field of spoken language processing is undergoing a shift from training custom-built, task-specific models toward using and optimizing *spoken language models* (SLMs) which act as universal speech processing systems. This trend is similar to the progression toward universal language models that has taken place in the field of (text) natural language processing. SLMs include both “pure” language models of speech—models of the distribution of tokenized speech sequences—and models that combine speech encoders with text language models, often including both spoken and written input or output. Work in this area is very diverse, with a range of terminology and evaluation settings. This paper aims to contribute an improved understanding of SLMs via a unifying literature survey of recent work in the context of the evolution of the field. Our survey categorizes the work in this area by model architecture, training, and evaluation choices, and describes some key challenges and directions for future work.

這個趨勢與文字自然語言處理 (text-based NLP) 朝向通用語言模型的發展路徑相似

1 Introduction

In the last few years, the field of natural language processing (NLP) has evolved from (1) training many task-specific models from scratch, to (2) combining pre-trained multi-purpose contextual representation models (such as BERT (Devlin et al., 2019)) with a small number of task-specific parameters, to (3) training generative *universal*, large language models (LLMs (Brown et al., 2020; OpenAI et al., 2024)¹) that perform arbitrary text tasks given natural language instructions (prompts) and can generalize to unseen domains and tasks (Wei et al., 2022a; Liu et al., 2023), and finally to (4) dialogue / chatbot systems that function as assistants and perform tasks while directly interacting with the user.

文字模型的發展:
專用->通用
->chatbot

The field of speech processing has been undergoing a similar evolution, although with some lag, and has mainly focussed on stages (1) and (2). The current state of the art (SOTA) for common specialized speech tasks—including automatic speech recognition (ASR), speech translation (ST), spoken language understanding (SLU) tasks, and speaker identification (SID)—involves combining a pre-trained self-supervised encoder model (Mohamed et al., 2022) with a task-specific prediction “head”. For some very common tasks—namely, ASR and ST—in relatively high-resource languages, large supervised models (Radford et al., 2023; Peng et al., 2023b) also have consistently good (if not SOTA) performance.

Recent work has begun to develop *spoken language models* (SLMs), analogous to text LLMs, which are in principle capable of performing arbitrary speech tasks given natural language instructions. However, the

*: Equal-contribution first authors listed in alphabetical order, +: Equal-contribution last authors listed in alphabetical order, #: Work was done before Haibin Wu joined Microsoft.

¹Throughout the paper, we use the terms LLMs and LMs interchangeably to refer to modern language models.

Table 1: Typology of text and spoken LMs. We use a loose notation here, where *speech* and *text* are to be interpreted in context; for example, $p(\text{text}|\text{text})$ in post-trained text LMs corresponds to modeling some desired output text given an input text instruction or prompt. “Post-training” refers to any form of instruction-tuning and/or preference-based optimization of the SLM. Please see the sections below for details and references for the example models.

Type of LM	Training Strategy	Model distribution	Examples
pure text LM	pre-training	$p(\text{text})$	GPT, Llama
pure text LM	post-training	$p(\text{text} \text{text})$	ChatGPT, Llama-Instruct
pure speech LM	pre-training	$p(\text{speech})$	GSLM, AudioLM, TWIST
pure speech LM	post-training	$p(\text{speech} \text{speech})$	Align-SLM
speech+text LM	pre-training	$p(\text{text}, \text{speech})$	SpiRit-LM, Moshi (pre-trained)
speech+text LM	post-training	$p(\text{text}, \text{speech} \text{text}, \text{speech})$	Moshi (post-trained), Mini-Omni
speech-aware text LM	post-training	$p(\text{text} \text{speech}, \text{text})$	SALMONN, Qwen-Audio-Chat

term “SLM” has not been standardized in the literature and has been used to refer to a wider range of model types than text language models. Several common classes of models have emerged, all referred to as SLMs: (1) pure SLMs: models of the distribution of speech, $p(\text{speech})$, typically trained on unlabeled tokenized speech data only with a next-token prediction objective (similarly to the pre-training phase in text LLMs) (Lakhotia et al., 2021); (2) speech+text SLMs: models of the joint distribution of speech and the corresponding text $p(\text{text}, \text{speech})$, typically trained using paired (speech, text transcription) data, which can be viewed as a direct extension of class (1) (Nguyen et al., 2025); and (3) speech-aware text LMs: models that combine text LLMs with pre-trained speech encoders to retain the instruction-following (IF) characteristics of the text LLMs and reason about the input speech or audio recordings (Tang et al., 2024; Chu et al., 2023). Approaches in category (3) model a conditional distribution, $p(\text{text}|\text{speech}, \text{text})$, of desired output text in response to the input speech and text instruction.

Following the text LLMs analogy, models in categories (1) and (2) can be viewed as analogues to pre-trained LLMs. Like text LLMs, these models can be post-trained—via instructions, preferences, or other means—to learn the distributions of desired output speech (for category (1)) or output speech+text (category (2)) given desired inputs (speech or speech+text, respectively). Models in category (3) typically start with fine-tuned LLMs, but involve some additional post-training to both align the speech and text representations and enable the model to perform new speech-specific tasks. Table 1 provides a typology of these categories along with example models from the recent literature. The Appendix provides additional models and a timeline of their development (Table 3 and Figure 6).

Although existing models have been applied to different tasks using different priors and modeling techniques, all of the three categories of SLMs mentioned above form steps on the way to universal speech processing systems. For the purposes of this paper, we define a universal speech processing system as a model that satisfies the following criteria:

1. It has both spoken input and spoken output with optional text input and/or output. The spoken input may serve as either an instruction or a context. 它同時具有語音輸入與語音輸出
2. It is intended to be “universal”; that is, it should in principle be able to address arbitrary spoken language tasks, including both traditional tasks and more complex reasoning about spoken data. 設計目標是成為**「通用型」模型**
3. It takes instructions or prompts in the form of natural language (either speech or text), and not, for example, only task specifiers (Radford et al., 2023) or soft prompts (Chang et al., 2024). 接受以自然語言形式給出的指令或提示（可以是語音或文字），而不是僅限於任務標識符或軟提示（soft prompts）

This is a functional definition, and does not restrict the form of the model. We also note that most of the models in the current literature, and therefore in this survey, do not satisfy all of these criteria; for example, many do not have speech as both input and output, and many are trained and evaluated on a fairly limited set of tasks. However, the models we include can all be seen as steps toward the same goal of eventually

developing SLMs that can serve as universal speech processing systems, in the same way that pre-trained and post-trained text LLMs have served as steps toward universal written language processing systems.

One may wonder whether a better path to universal speech processing is to combine speech recognition, text LLMs, and speech synthesis in series. This is indeed a strong baseline approach for many tasks (Huang et al., 2025; Yang et al., 2024b). However, some tasks require access to aspects of the audio signal beyond the word string, such as properties of the speaker, their emotional state, prosody, or other acoustic properties. In addition, even for tasks that are in principle addressable using the word string alone, an end-to-end SLM approach can avoid the compounding of errors and inefficiencies that can occur when combining ASR, text LLMs, and TTS systems in series.

The past few years have seen a major acceleration in work on SLMs. However, there have been limited efforts to perform an in-depth survey of the design and modeling choices made in this work. Additionally, these models are often evaluated on very different tasks and datasets, making it difficult to assess the relative performance of different approaches. As part of this survey, we collect and organize many of the existing evaluations (though standardized evaluation is still one of the remaining challenges in this research area). Lastly, as SLMs have been viewed and defined differently depending on the intended applications and modeling choices, we provide a unified formulation and terminology for describing SLMs (Section 2).

This survey is intended to serve as a snapshot of the current moment in the evolution of the field, providing a review of the recent literature and a unified definition of SLMs and their components. While new SLMs are being proposed at a steady pace, we hope that this survey of the research landscape will help readers more easily place new models in context. We aim to provide an improved understanding of the successes and remaining limitations of SLMs developed thus far, along the path to SLMs as universal speech processing systems.

Scope of this survey. Although the ultimate goal is task-independent models that can be instructed with natural language, many LM-inspired approaches thus far have been task-specific (for example, special cases of conditional $p(\text{text}|\text{speech})$ models for speech recognition and translation such as Whisper (Radford et al., 2023) and OWSM (Peng et al., 2023b), and conditional $p(\text{speech}|\text{text})$ models for text-to-speech synthesis such as VALL-E (Chen et al., 2025b) and VoiceBox (Le et al., 2024)), and some have been more general but rely on task tokens or soft prompts (e.g., Qwen-Audio (Chu et al., 2023) and SpeechPrompt (Chang et al., 2024) respectively). In this survey, we focus on models that are at least in principle task-independent (although they may have been tuned on a relatively small set of tasks) and that take natural language as input rather than task tokens or soft prompts. We will, however, discuss some of the important task-specific models as they relate to the more general models.

In addition, as discussed in Section 2, SLMs often comprise several components including a speech encoder, speech decoder, speech-text alignment module (when applicable), and sequence modeling component. In this survey, we provide relatively brief descriptions of speech encoding and decoding; these have been covered well in previous surveys (e.g., Wu et al. (2024a)), and our main focus is on the aspects that are specific to SLMs, such as sequence modeling and speech-text alignment.

Finally, music and general audio share many properties with speech. A number of language modeling approaches have been developed specifically for music and general audio (e.g., (Copet et al., 2023; Agostinelli et al., 2023)), and some SLM research combines speech and other audio (e.g., (Gong et al., 2023b)). In this survey we include non-speech audio only to the extent that it is used in the context of SLMs.

Related surveys. To place this paper in context, we note that several other recent surveys have addressed aspects of SLMs. Zhang et al. (2024a); Chen et al. (2024a) review multi-modal LLMs, including a limited subset of SLMs. By focusing on SLMs as our main subject, we survey a substantially larger set of models in greater detail and explore speech-specific issues. Latif et al. (2023) provide a survey on large audio models. This survey covers a much larger space of audio language models (including environmental sounds, music, and other audio) and therefore does not present SLMs in as much detail, nor include recent *instruction-following* (IF) models of the past year. Another recent survey (Wu et al., 2024a) provides an overview of neural audio codecs and codec-based (Section 3) models, which mainly focuses on speech tokenization rather

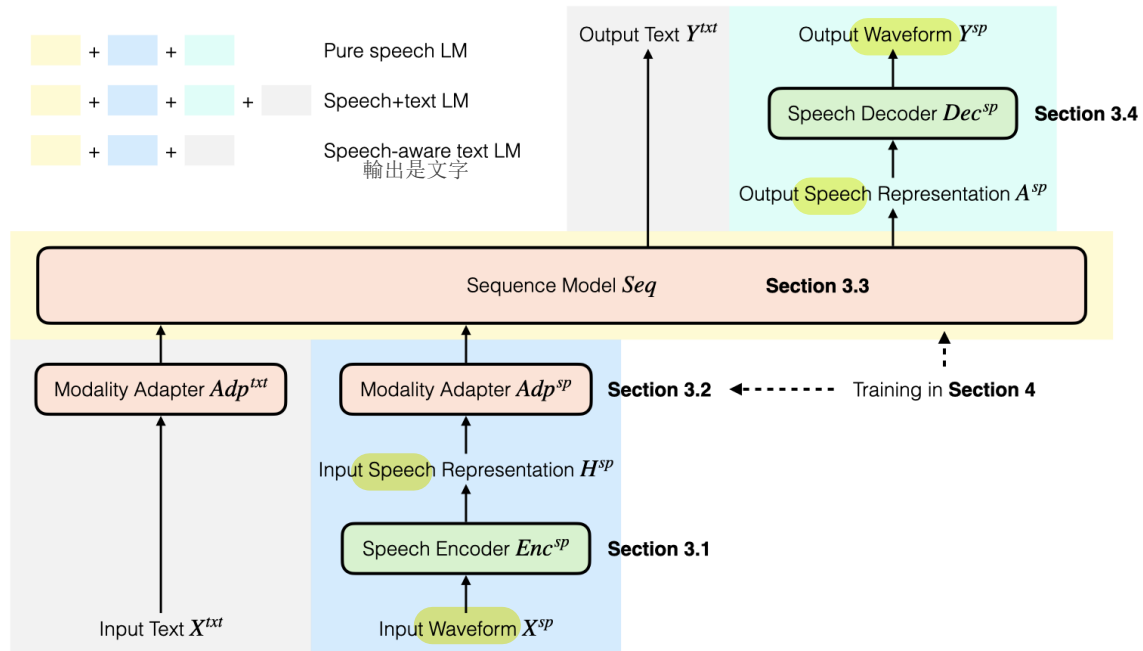


Figure 1: Overview of SLM architecture. See Sections 3 and 4 for more detailed descriptions of the components and training methods, respectively.

than the full range of SLMs. Guo et al. (2025) also survey tokenization methods. Ji et al. (2024) provide a survey of SLMs, but their main focus is on spoken dialogue systems. Another related SLM survey paper by Peng et al. (2024a) focuses on speech-aware text LMs but leaves out other types of SLMs. Finally, there is also concurrent work by Cui et al. (2024) that surveys a similar range of SLMs; we provide a complementary view of the landscape with a different historical focus and categorization of models.

Outline. In the remaining sections, we begin by outlining a general formulation of SLMs (Section 2), followed by a discussion of various design choices (Section 3), including speech encoders and decoders, modality adapters, and sequence modeling. We then describe the multiple optimization pipelines used for training SLMs (Section 4). In Section 5, we review and categorize some of the notable recent models. In Section 6, we discuss how SLMs have been adapted for dialogue (“full duplex”) mode. In Section 7, we present existing approaches for evaluating SLMs. Finally, we conclude by discussing the limitations of current approaches and provide some suggestions for future research (Section 8).

2 Overall architecture

We start by providing a *general* formulation (shown in Figure 1) of SLMs, which takes either/both speech and text as input and generates either/both speech and text (where at least one of the input or output includes speech). Although SLMs differ in many of their design choices, this unifying description subsumes all of the models we cover.²

Let $X^{\text{txt}} \in \mathcal{V}^*$ and $X^{\text{sp}} \in \mathbb{R}^*$ denote the input text and speech, respectively. That is,

任意長度N的文字序列 任意長度T的語音序列

$$X^{\text{txt}} = \{t_1, t_2, \dots, t_N\}$$

²We do make some assumptions, for example, that the sequence model is autoregressive, which in principle need not hold but in practice do hold for current models.



is a sequence of text tokens t_i of arbitrary length N drawn from a vocabulary \mathcal{V} . The speech input is a waveform, that is a sequence

$$X^{\text{sp}} = \{s_1, s_2, \dots, s_T\} \quad \text{通常語音長度} \gg \text{文字長度} \quad \rightarrow \text{Tokenization的議題}$$

of real-valued samples of arbitrary length T , where typically $T \gg N$. The speech waveform is long and complex, and is therefore typically further processed before being modeled by the sequence model. We first map X^{sp} using the speech encoder $\text{Enc}^{\text{sp}}()$ to a sequence of speech representations H^{sp} :

把原始資料（例如語音、圖像、文字）轉換成模型可以理解、計算、處理的向量形式

$$H^{\text{sp}} = \text{Enc}^{\text{sp}}(X^{\text{sp}}). \quad (1)$$

H^{sp} 可能是連續的向量或是被轉換成離散的token

The resulting speech representations, $H^{\text{sp}} = \{h_1, h_2, \dots, h_L\}$, can take one of two forms. In one case, H^{sp} is a sequence of continuous vectors $h_l \in \mathbb{R}^d$. Alternatively, they can be sequences of discrete tokens $h_l \in \mathcal{D}$ drawn from a learned codebook \mathcal{D} .

The speech representations H^{sp} are further transformed through a modality adapter

$$\text{Adp}^{\text{sp}}(H^{\text{sp}}) \in \mathbb{R}^{d' \times L'},$$

d' : 語音特徵的新維度

L' : sequence length可能比原本少，用來對齊文字序列長度或簡化輸入

where typically $L' \leq L$. This transformation is intended to improve the alignment of the speech representations with the sequence model. In addition, the inherent length disparity between speech and text sequences can be addressed at this stage by applying temporal adjustments in $\text{Adp}^{\text{sp}}()$. The text token sequence X^{txt} is also transformed into a sequence of vector representations, with the same dimensionality d' , via an adapter

$$\text{Adp}^{\text{txt}}(X^{\text{txt}}) \in \mathbb{R}^{d' \times N'}.$$

The text adapter is typically simply an embedding table, and therefore $N' = N$.³ After the adapters, therefore, the text and speech are mapped to the same representation space with dimension d' , but the text and speech representation sequences are still not necessarily the same length (i.e., L' and N' need not be identical).

Next, given the adapted input text and/or speech representations, a sequence model, denoted $\text{Seq}()$, generates outputs, typically in an autoregressive manner. Here, we assume that each invocation of $\text{Seq}()$ corresponds to one generation step. For the model types described in Table 1, the formulations are as follows:

Pure Speech LMs The output of $\text{Seq}()$ is a speech representation

$$h' = \text{Seq}(\text{Adp}^{\text{sp}}(H^{\text{sp}})).$$

The input to $\text{Seq}()$ is a sequence of representations of the generated speech so far, with shape $d' \times L'$ and the output speech representation h' is either a continuous vector (in \mathbb{R}^d) or a discrete token (in \mathcal{D}). Following the usual autoregressive generation formulation, h' is then appended to H^{sp} , increasing its length by one, $\text{Seq}()$ then generates the next output, and so on. All of the generated outputs h' together form a sequence $A^{\text{sp}} = \{h'_1, h'_2, \dots, h'_{L''}\}$ with sequence length L'' . Finally, the speech decoder Dec^{sp} converts the output speech representations A^{sp} to a waveform Y^{sp} :

最終用decoder轉回音訊

$$Y^{\text{sp}} = \text{Dec}^{\text{sp}}(A^{\text{sp}}). \quad (2)$$

Speech-aware Text LMs In this type of model, the sequence model $\text{Seq}()$ generates a text token t' according to

$$t' = \text{Seq}(\text{Adp}^{\text{sp}}(H^{\text{sp}}), \text{Adp}^{\text{txt}}(X^{\text{txt}}))$$

也是一種文字接龍

Here, the input to $\text{Seq}()$ is a concatenation of the adapted speech representation $\text{Adp}^{\text{sp}}(H^{\text{sp}})$ and the text representation $\text{Adp}^{\text{txt}}(X^{\text{txt}})$, forming a tensor with shape $d' \times (L' + N')$. The output is a text token $t' \in \mathcal{V}$. This generated token is then appended to X^{txt} , and the process is repeated to generate subsequent tokens. Finally, the sequence of all generated tokens forms the output text sequence Y^{sp} .

³We are not aware of any SLMs that use any text adapters besides the typical embedding table, nor models where $N' \neq N$, but we allow this possibility in the definition for maximal generality.

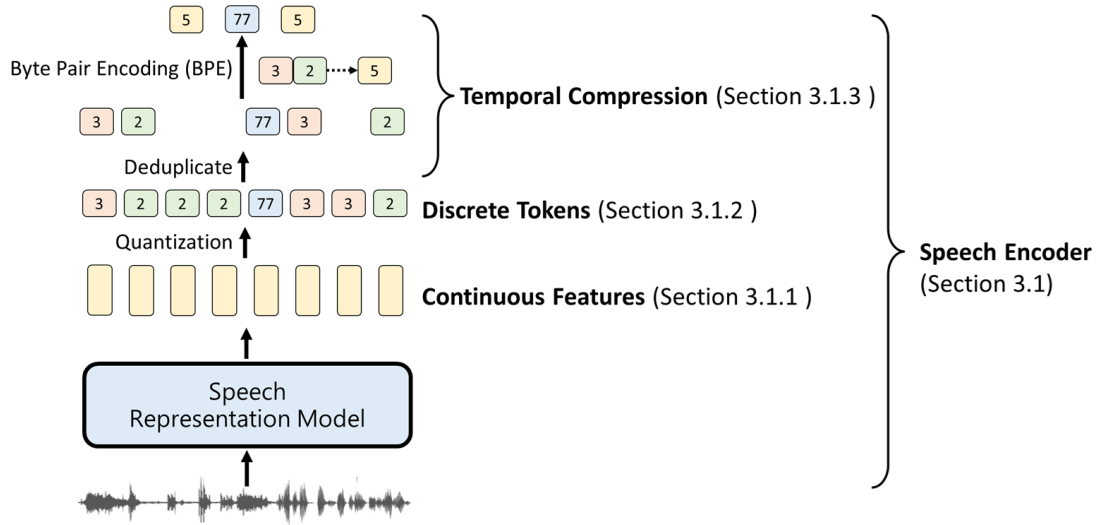


Figure 2: A general pipeline for speech encoders. Note that different encoders use different components of the pipeline. See Section 3.1 for more details.

Speech+Text LMs In this more complex scenario, speech and text are modeled jointly:

$$h', t' = \text{Seq}(\text{Adp}^{\text{sp}}(H^{\text{sp}}), \text{Adp}^{\text{txt}}(X^{\text{txt}}))$$

Here both the inputs and outputs are “hybrid” representations consisting of a combination of speech and text representations. There are multiple approaches for generating such hybrid speech+text representations, which will be discussed in Section 3.3.2.

The encoder, decoder, and sequence model $\text{Seq}()$ are often pre-trained separately. The modality adapter is typically trained from scratch, sometimes along with fine-tuning of the sequence model. The encoder and decoder parameters are usually fixed after pre-training. Section 3 provides more detailed information about the model components, while Section 4 describes typical training methods.

模態轉接器 (modality adapter) 則通常是從頭開始訓練，有時會與序列模型一同進行微調 (fine-tuning)。在預訓練完成之後，編碼器與解碼器的參數通常會固定

3 SLM components

3.1 Speech Encoder

In text LMs, text is typically tokenized into subword units, generated through methods like byte pair encoding (BPE) (Gage, 1994), and these tokens are then represented as vector embeddings. In contrast, speech is a continuous waveform with no obvious tokenization and no separation between linguistic information and other acoustic aspects. The task of the speech encoder, shown in Figure 2, is to extract meaningful representations from the waveform.

The first part of the encoder is a speech representation model, which transforms the speech signal into continuous features (vectors). These continuous features can either be used directly as input to the speech modality adapter $\text{Adp}^{\text{sp}}(\cdot)$ or quantized into discrete units (using, e.g., k -means or VQ-VAE (van den Oord et al., 2017)). In general, pure speech LMs and speech+text LMs tend to use discrete speech tokens, whereas speech-aware text LMs tend to use continuous representations (with some exceptions; e.g., Flow-Omni (Yuan et al., 2024) is a speech+text LM that generates continuous speech representations, and DiscreteSLU (Shon et al., 2024) is a speech-aware text LM that uses discrete speech token input). Finally, temporal compression (e.g., deduplication or BPE) is optionally applied to reduce the sequence length.

語音表示模型：
會將語音訊號轉換為
連續特徵向量
(continuous
features)

3.1.1 Continuous Features

To extract informative representations from raw waveforms, a speech representation model—either a learned encoder or a digital signal processing (DSP) feature extractor—converts speech into continuous features. These continuous features may include:

1. Traditional spectrogram features, such as mel filter bank features (Huang et al., 2001).
2. Hidden representations from self-supervised learning-based (SSL) speech encoders, such as wav2vec 2.0 (Baevski et al., 2020), HuBERT (Hsu et al., 2021), or WavLM (Chen et al., 2022).
3. Hidden representations from supervised pre-trained models, such as Whisper (Radford et al., 2023) or USM (Zhang et al., 2023b).
4. Hidden representations from neural audio codec models, such as SoundStream (Zeghidour et al., 2022) or EnCodec (Défossez et al., 2023).

3.1.2 Discrete Tokens

We divide discrete speech tokens into two main categories, “phonetic tokens” and “audio codec tokens”, based on how they are learned and their perceived properties.⁴

Phonetic Tokens Phonetic tokens are most commonly derived by quantizing self-supervised speech representations (Mohamed et al., 2022) or, occasionally, supervised encoders from a pre-trained ASR model (as in Cosyvoice (Du et al., 2024), GLM-4-voice (Zeng et al., 2024a), and WhisperSpeech (WhisperSpeech, 2024)). These quantized representations have been shown to capture linguistic information and have strong correlation with phonetic units (Sicherman & Adi, 2023). The pioneering pure speech LM, GSLM (Lakhotia et al., 2021), treated these tokens as “pseudo text,” allowing the model to function similarly to text-based LMs and to generate phonetically coherent speech. These tokens, learned directly from raw speech, formed a foundation for subsequent research, often described as “textless NLP” (Polyak et al., 2021; Lakhotia et al., 2021; Kharitonov et al., 2022b; Nguyen et al., 2023; Hassid et al., 2023).

第12講P. 17

The process of deriving phonetic tokens involves several design decisions, such as selecting a pre-trained representation model, which layer’s representations to quantize, and the number of clusters. Different layers in SSL speech models tend to encode different types of information (Hsu et al., 2021; Mousavi et al., 2024b; Pasad et al., 2023). Representations from layers rich in phonetic information are typically chosen for extracting phonetic tokens. The number of clusters is also a key hyperparameter, and is often optimized based on downstream task performance, or the ABX task (Schatz et al., 2013) which is a good proxy for downstream performance. Too few clusters may result in a loss of fine-grained phonetic information, while too many clusters may risk encoding undesirable speaker-specific features (Lakhotia et al., 2021; Kharitonov et al., 2022a).

Audio Codec Tokens In contrast to phonetic tokens, audio codec tokens are intended to capture more detailed acoustic characteristics. These tokens are derived from neural codec models, which were originally designed for audio compression and therefore for faithful reconstruction of audio from its encoding (Zeghidour et al., 2022; Défossez et al., 2023). Although audio codecs were initially intended for compression, their intermediate discrete representations have proven valuable as tokens in SLMs (Borsos et al., 2023a).

Neural codecs comprise three components: an encoder that converts raw audio into frame-level features, a vector quantization (VQ) module that converts these features into discrete tokens, and a decoder that reconstructs audio from the codec tokens. Typically, the main encoder-decoder backbone is a convolution-dominant architecture (e.g., Encodec (Défossez et al., 2023)), but some recent work has introduced transformers as intermediate layers (e.g., Mimi (Défossez et al., 2024)) or transformer-only architectures to reduce computation (e.g. TS3-Codec (Wu et al., 2024d)).

⁴In some literature, phonetic tokens are referred to as “semantic tokens” (Borsos et al., 2023a). However, it’s important to clarify that “semantic” in this context does not imply the traditional linguistic meaning, as these tokens are more akin to phonetic units (Sicherman & Adi, 2023) and do not typically carry semantic content. Therefore, we use the term “phonetic tokens” throughout this paper.

第12講P. 19

For the VQ module, a commonly used approach is residual vector quantization (RVQ) (Gray, 1984; Zeghidour et al., 2022), which generates several discrete tokens for each time step, corresponding to multiple (hierarchical) levels of granularity (with the first level encoding most of the information, and the remaining levels encoding residual information). Decoding such multi-codebook speech tokens typically requires additional design considerations for the SLM architecture and decoding strategy (Borsos et al., 2023a; Chen et al., 2025b; Yang et al., 2024a). Alternatively, several single-codebook quantization codecs (Wu et al., 2024d; Xin et al., 2024) have been developed to simplify decoding in SLMs. Section 3.3 describes decoding strategies in detail.

Comparison of token types Early SLMs, such as GSLM, typically used phonetic tokens, which reduce speaker-specific information (Polyak et al., 2021). This reduction allows language models to focus primarily on spoken content, making them more effective for understanding tasks like detecting words vs. nonwords or syntactic correctness (Lakhotia et al., 2021; Hassid et al., 2023) and for applications such as speech-to-speech translation (Lee et al., 2022a;b). In contrast, audio codec tokens are frequently used in tasks where preserving speaker identity and acoustic details is crucial (Chen et al., 2025b).

The development of speech tokenization methods has attracted increasing attention. In SLMs, two important factors are the token bitrate, which impacts efficiency, and the token quality, which relates to generation quality and suitability for downstream tasks. Several benchmarks have been established to evaluate different types of tokens (Shi et al., 2024; Wu et al., 2024c; Mousavi et al., 2024a; Wu et al., 2024b). Codec-SUPERB (Wu et al., 2024c), the first neural audio codec benchmark, evaluates the quality of resynthesized audio, using subjective metrics and pre-trained downstream models for comparison. DASB (Mousavi et al., 2024a) evaluates tokenization methods by using the extracted tokens for various downstream tasks. ESPnet-Codec (Shi et al., 2024) is an open-source framework that functions as both a toolkit for neural codec training and a platform for evaluation in a controlled setting. Designing efficient and effective tokens for SLMs remains an active area of research (Guo et al., 2025).

3.1.3 Temporal Compression

Quantized speech feature sequences are often very long due to the high frame rate of speech signals. To mitigate the challenges this poses for language modeling, several techniques are typically applied within the encoder to reduce the sequence length. One such technique is “deduplication”, which merges consecutive identical tokens into a single token (as shown in Figure 2). However, this approach loses information about the duration of individual tokens. To address this issue, some approaches (Nguyen et al., 2023) use multi-stream modeling to capture token duration in a separate stream, while others introduce duration prediction when generating output speech (Kreuk et al., 2022; Maimon & Adi, 2023; Lee et al., 2022a). Additionally, byte pair encoding (BPE) (Gage, 1994) is also sometimes applied to the discrete token sequences (Wu et al., 2023a; Shen et al., 2024) to capture recurring patterns.

3.2 Speech Modality Adapter

許多SLMs（尤其 speech-aware text LMs），語音編碼器與序列模型最初是分開開發的，後來再整合。因此，有必要將語音編碼器的輸出對齊到序列模型所期望的輸入格式，而這正是modality adapter的角色。模態轉換器通常會在下游任務上訓練；如果是speech+text LMs，則可能會在預訓練階段一起訓練

In many SLMs (especially speech-aware text LMs), the speech encoder (Section 3.1) and the sequence model (Section 3.3) are initially developed separately and then combined. It is therefore necessary to somehow align the output of the speech encoder with the expectations of the sequence model, and this is the role of the modality adapter. The modality adapter is typically trained on downstream tasks or, in the case of speech+text LMs, as part of pre-training (see Section 4 for more details on training).

If the output sequence of the speech encoder is not very long, the modality adapter can be as simple as a linear layer, which transforms the encoder output into the embedding space of the sequence model. This typically occurs when the encoder produces discrete tokens with temporal compression. On the other hand, if the output of the encoder is too long, the adapter is typically also responsible for shortening the sequence. This usually happens when the encoder produces continuous representations without discretization or temporal compression. Shortening the input sequence simplifies both training and inference for the sequence model. Furthermore, since the sequence model often processes both text and speech inputs, it is helpful for the token rate of the speech sequence to roughly match that of the text sequence.

根據輸出序列長度，有不同的設計

Common adapters include:

Linear Transformation / Vocabulary Expansion A straightforward way to integrate speech into a pre-trained sequence model is by applying a linear transformation to the speech representation H^{sp} . In this approach, the adapter $\text{Adp}^{\text{sp}}()$ is defined as a linear transformation. This method is commonly used when the speech encoder produces discrete token outputs. This approach can be interpreted as vocabulary expansion, where the speech tokens are treated as additional tokens in the sequence model’s vocabulary. Their embeddings are then learned during subsequent task-oriented training. For example, in SpeechGPT (Zhang et al., 2023a), the vocabulary of the sequence model LLaMA (Touvron et al., 2023a) is expanded to include HuBERT-based phonetic tokens. Similarly, in Mini-Omni2 (Xie & Wu, 2024b), the sequence model Qwen (Bai et al., 2023) incorporates eight layers of acoustic codec tokens alongside standard text tokens.

P.5

CNN with strides: Convolution layers with strides reduce the sequence length while preserving temporal information (Lu et al., 2024a), which is essential for tasks that require such information, like ASR (Wu et al., 2023b). A special case of this adapter type is pooling layers with strides (Chu et al., 2023).

在這裡輸

Connectionist Temporal Classification (CTC)-based compression: This method compresses H^{sp} (Eq. 1) according to the posterior distribution from a CTC-based speech recognizer (Gaido et al., 2021). CTC (Graves et al., 2006), a commonly used approach for ASR, assigns each time step a probability distribution over a set of label tokens, including a blank (“none of the above”) symbol. The time steps with high non-blank probabilities indicate segments that are likely to carry important linguistic information. CTC compression aggregates the frame-level labels, specifically by merging repeated non-blank labels and removing blanks. This approach produces a compressed representation intended to retain the relevant content of the original sequence while significantly reducing its length (Wu et al., 2023b; Tsunoo et al., 2024).

Q-Former: The Q-Former (Li et al., 2023) is an adapter that produces a fixed-length representation by encoding a speech representation sequence of arbitrary length into M embedding vectors, where M is a hyperparameter (Lu et al., 2024b).

Let the input speech representation sequence be:

$$X = \{x_1, x_2, \dots, x_{L'}\}, \quad x_i \in \mathbb{R}^{d'}, \quad (3)$$

where L' is the sequence length and d' is the dimension of the embeddings.

To achieve a fixed-length representation, Q-Former introduces M trainable query embeddings:

$$Q = \{q_1, q_2, \dots, q_M\}, \quad q_i \in \mathbb{R}^{d'}. \quad (4)$$

These queries interact with X via a cross-attention mechanism:

$$\text{Attn}(Q, X) = \text{softmax}\left(\frac{QW_Q(XW_K)^T}{\sqrt{d'}}\right)XW_V, \quad (5)$$

where W_Q , W_K , and W_V are learnable projection matrices. The result is a sequence of M embeddings.

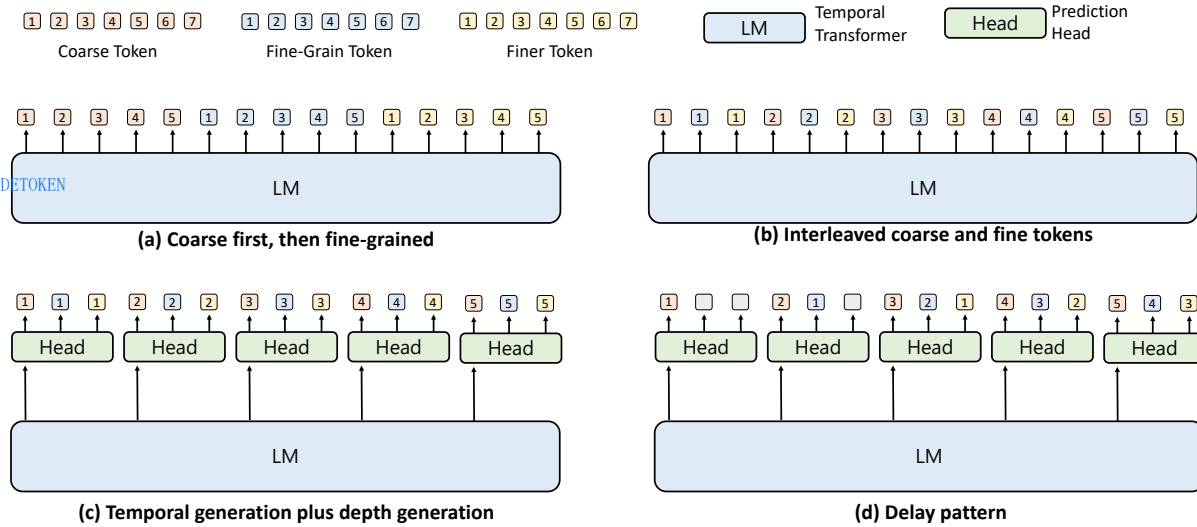
In some approaches, instead of directly encoding the entire utterance into M vectors, a *window-level Q-Former* is applied (Yu et al., 2024; Pan et al., 2024; Tang et al., 2024) to retain temporal information. In the window-level Q-Former, the input embedding sequence is segmented, and the Q-Former is applied to each segment.

Lu et al. (2024a) compare the Q-Former with CNN-based modality adapters in a speech-aware text LM, finding that the Q-Former produces better performance on the Dynamic-SUPERB benchmark (Huang et al., 2024) (see Section 7 for more on this and other SLM benchmarks).

AlignFormer: AlignFormer (Fan et al., 2024) combines a CTC compressor (Gaido et al., 2021) with a Q-Former (Li et al., 2023). When the LLM backbone is frozen, Fan et al. (2024) find that AlignFormer enables zero-shot instruction-following capabilities for speech QA tasks using only ASR training data. Additionally, AlignFormer surpasses Q-Former in instruction-following performance across multiple datasets and tasks, including SQA and speech translation.

第12講P. 23
需要等待前面的才能D
缺點：無法streaming

第12講P. 29



第12講P. 30

Figure 3: Hierarchical generation strategies (see Section 3.3.1).

Others: Several other methods have been developed for the modality adapter. For example, LTU-AS uses time and layer-wise transformers (TLTR) (Gong et al., 2023a), WavLLM (Hu et al., 2024b) uses a bottleneck adapter (Houlsby et al., 2019), and other approaches use multi-layer perceptron (MLP) adapters (Fang et al., 2025; Microsoft et al., 2025).

3.3 Sequence Model

In principle, generating speech tokens is similar to generating text tokens. However, there are some important differences. Unlike text tokens, audio tokens may include a mix of coarse and fine tokens — for example, phonetic tokens as coarse tokens and audio codec tokens as fine tokens. For TTS models, using phonetic tokens as an intermediate representation rather than directly mapping text to audio codec tokens provides better supervision and reduces modelling complexity. For example, SPEAR-TTS (Kharitonov et al., 2023), which uses phonetic tokens, has similar performance to VALL-E (Chen et al., 2025b) while requiring significantly less parallel text-to-speech training data. The same concept applies to SLMs, as described in Section 3.3.1 below. In addition to generating multiple speech token types, some SLMs combine text and speech token generation to enhance the linguistic quality of the generated speech, as described in Section 3.3.2.

第12講P. 20

3.3.1 Hierarchical Token Generation

There are several decoding strategies for hierarchical modeling of multiple types of tokens of different granularity, shown in Figure 3. Phonetic tokens and first-layer codec tokens (e.g. Mimi (Défossez et al., 2024) and SpeechTokenizer (Zhang et al., 2024b)) can be regarded as coarse tokens in Figure 3, while the remaining layers of codec tokens can be regarded as fine-grained and finer tokens.⁵ We categorize common decoding strategies into the following types:

Coarse first, then fine-grained (Figure 3 (a)): In this approach, coarse tokens for all time steps are generated first, followed by fine-grained tokens conditioned on the coarse tokens, and finally the finer tokens. For example, AudioLM (Borsos et al., 2023a) uses such a three-stage autoregressive process (which can be modeled by different LMs), first predicting phonetic tokens (the coarse tokens in Figure 3 (a)), then first-layer audio codec tokens (the fine-grained tokens in Figure 3 (a)) conditioned on the phonetic tokens, and finally the remaining-layer codec tokens (the finer audio codec tokens in Figure 3 (a)). SoundStorm (Borsos

⁵In Figure 3 and in running examples, we assume 3 granularity levels. In practice, a sequence model can include a smaller or larger number of levels, and this number is also an important design choice.

et al., 2023b) replaces the second and third stages of AudioLM with a non-autoregressive prediction approach inspired by MaskGIT (Chang et al., 2022), which begins with masked tokens and then non-autoregressively predicts a portion of them in rounds based on confidence scores. Similarly, VALL-E (Chen et al., 2025b) uses an autoregressive model to predict the first-layer audio codec tokens based on the corresponding phoneme transcriptions, and then uses a non-autoregressive model to generate the remaining codec token sequences.

Interleaved coarse and fine tokens (Figure 3 (b)): In this decoding strategy, the three types of tokens—coarse, fine-grained, and finer—are aligned along the time axis and interleaved for generation. When predicting tokens for time step t , the coarse, fine-grained, and finer tokens corresponding to t are generated sequentially. SpiRit-LM (Nguyen et al., 2025) applies this interleaved structure (interleaved phonetic, pitch and style tokens) to enhance speech expressiveness, and it also includes text tokens in addition to the multiple types of speech tokens.

Temporal generation plus depth generation (Figure 3 (c)): This strategy employs a large Transformer LM (indicated in blue) to model inter-frame information along the time axis and a small Transformer head (indicated in green) to capture intra-frame correlations. The small Transformer head predicts multi-layer tokens—where fine-grained and finer tokens correspond to different layers of audio codec tokens—within the same time step. For example, UniAudio (Yang et al., 2024a) adopts this approach, inspired by RQ-Transformer (Lee et al., 2022c) and the MegaByte Transformer model (Yu et al., 2023). Moshi (Défossez et al., 2024) also adopts a similar strategy as UniAudio.

Delay pattern (Figure 3 (d)): In this approach, there are time delays (of 1 or more time steps, depending on the design) between coarse and fine tokens. Such time delays allow the model to utilize look-ahead, that is to use future coarse tokens as conditions when generating fine tokens. Figure 3 (d) shows an example where the delay is 1. The large blue LM in Figure 3 (d) autoregressively outputs the temporal embeddings and feeds them into the green prediction head, which then predicts the coarse token and the delayed fine token in parallel. For example, pGSLM (Kharitonov et al., 2022b) introduces a delay between the phonetic tokens and the “expressive” (pitch and duration) tokens.

Using multiple stages of prediction can achieve both high audio quality and long-term consistency (Borsos et al., 2023a). However, this approach makes the decoding strategy complex and increases latency, making it less suitable for real-time applications like spoken dialogue generation. To address this drawback, there is a growing literature on tokenization methods that reduce the number of tokenization layers. One example is to combine the generation of phonetic and audio codec tokens into a single tokenizer by distilling the information from phonetic tokens directly into the audio codec (Défossez et al., 2024; Zhang et al., 2024b).

3.3.2 Text and speech hybrid generation

Some sequence generation models use text tokens as the first-layer tokens, followed by speech tokens (either phonetic or audio codec tokens). This approach allows the SLM to leverage the knowledge from a pretrained text LM, which is typically trained on much larger-scale data than the available speech data, to improve the factuality and linguistic quality of generated speech. As shown by Défossez et al. (2024), this approach can improve language modeling quality, as measured by negative log-likelihood scores measured by an external text LM.⁶ Another advantage to this approach is that during the process of developing the SLM, it is straightforward to separately evaluate both the correctness of the output content and the speech generation quality.

Text and speech tokens operate on different scales: speech has a fixed sampling rate (unless tokenized into tokens of different durations), whereas text tokens do not. A text token sequence is also usually shorter than the corresponding speech token sequence. Hybrid decoding needs to address the issue of differing lengths and, ideally, also to achieve temporal alignment (synchronization). There are four main types of text-speech hybrid generation used in recent SLMs, shown in Figure 4. We use four representative models to illustrate the ideas.

⁶In this case, LiteLlama-460M-1T, <https://huggingface.co/ahxt/LiteLlama-460M-1T>



Figure 4: Text and speech hybrid generation, exemplified by four representative models (a)-(d) (see Section 3.3.2). Green and yellow boxes represent text and audio tokens, respectively.

The first type (Figure 4 (a)) addresses the mismatch in sequence lengths by adding padding tokens after the text token sequence. In this setup, the text sequence ends first, and the generated text then guides the generation of the speech token sequence, making the process similar to text-to-speech (TTS). A representative example is Mini-Omni (Xie & Wu, 2024a), which adopts a delayed decoding approach (Copet et al., 2023).

The second type (Figure 4 (b)) involves adding fixed padding tokens after each text token to extend the text token sequence. Padding tokens are also added between speech tokens so that both sequences have the same length. An example of this approach is LLaMA-Omni (Fang et al., 2025): After predicting the text tokens (padded with fixed-length padding tokens), it then predicts the phonetic token sequence based on the padded text (where the phonetic token sequence is shorter than the padded text token sequence), and finally the CTC loss is applied to the phonetic token sequence to guide training.

The third type (Figure 4 (c)) dynamically adjusts the number of padding tokens between text tokens. The generative model learns to insert these dynamic padding tokens to make the text and speech sequences the same length. During training, time-aligned text-speech paired data is used to construct padded text sequences that match the lengths of the speech sequences. Moshi (Défossez et al., 2024) is an example of this model type.

The fourth type (Figure 4 (d)) interleaves speech and text tokens in a single sequence, with text-speech token alignments derived in advance. SpiRit-LM (Nguyen et al., 2025) is an example of this approach, using time-aligned text-speech paired data for training. GLM-4-Voice (Zeng et al., 2024a;b) uses a pre-trained text-to-token model to generate audio tokens from text, and interleaves the generated audio with text tokens.

要把語音表徵轉回音訊

3.4 Speech Decoder

The speech decoder converts speech representations—whether continuous features, phonetic tokens, or audio codec tokens—back into waveforms. The speech decoder can take various forms:

離散

1. Vocoder (Kong et al., 2020) for continuous features, similar to those used in traditional synthesis systems. For instance, in Spectron (Nachmani et al., 2024), a generated mel spectrogram is synthesized into audio using the WavFit vocoder (Koizumi et al., 2023).

第12講P.61 Align-SLE論文用此

1. 對於連續特徵 (continuous features)，可以使用 Vocoder來進行語音合成，這與傳統的語音合成系統類似。

2. 對於語音單元 (phonetic tokens)，可以使用基於 HiFi-GAN 的 unit-based vocoder。這類 vocoder 以語音單元作為輸入，並可選擇性地結合額外資訊以提升合成品質。例如，當語音單元經過 deduplication 時，vocoder 通常會加入長度建模模組。

3. 對於 audio codec tokens，可使用對應的 pre-trained audio neural codec decoder 直接將這些 token 解碼為波形 (waveform)，無需額外訓練。

2. Unit-based vocoder (Polyak et al., 2021) based on HiFi-GAN (Kong et al., 2020) for phonetic tokens. These vocoders take phonetic tokens as inputs and optionally combine them with additional information to improve synthesis quality. For example, when phonetic tokens are deduplicated, a duration modeling module is often included in the vocoder (Lakhotia et al., 2021).
3. Codec decoder (Guo et al., 2025). When the SLM generates audio codec tokens, these tokens can be input directly into the corresponding pre-trained audio neural codec decoder (without additional training) to get the waveform.

4 Training Strategies

We divide the training phases of SLMs into *pre-training* and *post-training*, analogously (but not identically) to the typical division in text LLM training. For text LMs, pre-training refers to training a model of $p(\text{text})$ with a next-token prediction objective, whereas post-training is task-oriented and directly facilitates alignment and improved performance on tasks. Since many SLMs start with a post-trained text LLM, we consider this to be a part of SLM pre-training. More precisely, we define *pre-training* and *post-training* of SLMs as follows:

對於文字語言模型 (text LMs) 來說, pre-training 是指使用「下個詞預測目標」來訓練一個模型 $p(\text{text})$; 而 post-training 則是針對任務進行的訓練, 目的在於讓模型更好地對齊任務需求並提升任務表現。

Pre-training: We define pre-training as any training strategy that *does not have the explicit goal of enabling the model to perform a wide variety of downstream tasks*. In particular, pre-training does not directly aim to make a model a universal speech processing system (but, as described below, it can include training for specific tasks). In the context of SLMs, which handle multiple modalities, pre-training often involves a multi-stage process. This may include initial next-token-prediction training on text, followed by continual training on speech continuation tasks using large unlabeled speech corpora. Continual training may enhance the capabilities of SLMs for specific tasks, but the resulting pre-trained SLMs remain limited in scope and are still far from being universal models. Text LLM post-training before adding the speech modality is also considered pre-training of the SLM, since it does not target universal *speech* processing.

由於許多SLMs一開始就使用一個已經經過post-training的text LLM作為基礎, 所以我們將這個後訓練過程視為 SLM 預訓練的一部分。

Post-training: Post-training refers to training strategies *focused on variety of downstream tasks, with a goal of producing a more or less universal speech processing system*. The tasks can be either predefined with specific task specifiers or dynamically defined through text or spoken language prompts. Instruction tuning (Tang et al., 2024; Pan et al., 2024; Shu et al., 2023) and preference optimization (Lin et al., 2024b) are examples of post-training.

第6講P. 12

4.1 Pre-training strategies

Here we focus on speech-specific pre-training strategies; that is, we do not discuss the training of any text LM that is included in SLM pre-training.

4.1.1 Generative pre-training

Pure speech pre-training $p(\text{speech})$ Pure speech pre-training refers to training a model of $p(\text{speech})$ from unlabeled speech (a pure speech LM in Table 1), typically by tokenizing the speech and modeling the token sequences with an autoregressive model trained to perform *next speech token prediction*. The success of SSL speech representation models Mohamed et al. (2022) facilitated the development of pure speech LMs, because SSL models provided high-quality representations that could easily be quantized to produce tokenized speech. Pure speech LMs, such as GSLM (Lakhotia et al., 2021), typically use discrete phonetic speech tokens as the basic language modeling units, but *non-phonetic information*—such as duration and pitch—can also be incorporated (Kharitonov et al., 2022b; Nguyen et al., 2023) to enhance prosody modeling. Section 5.1 expands on pure speech LMs.

Joint speech and text pre-training $p(\text{text}, \text{speech})$ This approach refers to pre-training SLMs on aligned speech and text to model $p(\text{text}, \text{speech})$ (speech+text LM in Table 1) (Chou et al., 2023; Nguyen et al., 2025; Défossez et al., 2024). The speech and text sequences may be interleaved (Chou et al., 2023; Nguyen et al., 2025) or treated as dual channels (Défossez et al., 2024), as detailed in Section 3.3.2.

交錯處理

Continual pre-training following text pre-training $p(\text{text})$ Continual pre-training (Ke et al., 2023) refers to the process of further training a pre-trained model on additional data that is domain- or modality-specific, but still not task-oriented. This intermediate step allows the model to adapt to new domains or datasets while (hopefully) retaining its general-purpose capabilities (Tang et al., 2024; Chu et al., 2023; Hassid et al., 2023). A common example of continual pre-training for SLMs is to start with a text LLM pre-trained with a next token prediction objective (a pure text LM in Table 1) and then further pre-train it with speech data to improve its performance on the speech continuation task (Hassid et al., 2023).

EX. 續寫

4.1.2 Conditional pre-training $p(\text{text}|\text{speech})$

While generative pre-training is the most common form of pre-training, it is also in principle possible to initialize SLM training from a *conditional* model. For example, the pre-trained conditional model could be an encoder-decoder speech recognizer Chan et al. (2016), or a joint speech recognition and translation model like Whisper Radford et al. (2023) or OWSM Peng et al. (2023b; 2024c) that uses task specifier tokens to indicate the desired task. Such models can be trained on massive amounts of transcribed speech, and therefore have already learned to align the speech and text modalities. Although these models are trained for very specific tasks, they have shown promise as an initialization for instruction-following models that perform a wider range of understanding tasks (Lai et al., 2023; Arora et al., 2024).

4.1.3 Aligning speech and text modalities

For SLMs that combine pre-trained text LMs and speech encoders, another pre-training strategy is to align the speech and text modalities in a task-independent way.

Implicit alignment Speech and text modalities can be implicitly aligned through techniques such as the “modal-invariance trick” (Fathullah et al., 2024) or behavior alignment (Wang et al., 2023a). The idea is that the model should produce identical responses regardless of the input modality, provided the input conveys the same meaning. This approach often utilizes ASR datasets. The text transcript is input to a text LLM to generate a text response, while the corresponding speech recording is input into the SLM, which is trained to generate the same text response. Another idea found to be useful for implicit alignment is training spoken LLMs for audio captioning, where a spoken LLM takes audio as input and outputs its description. It has been observed that training a spoken LLM solely through audio captioning can generalize to tasks it has never seen during training (Lu et al., 2024a;b).

Explicit alignment Speech and text modalities can also be explicitly aligned by matching speech features to corresponding text embeddings, via optimization of appropriate distance/similarity measures. For example, Wav2Prompt (Deng et al., 2024) and DiVA (Held et al., 2024) align modalities by minimizing the L_2 distance between speech features and the token embeddings of their transcripts in a text LLM while keeping the text embeddings fixed.

4.2 Post-training strategies

The pre-training of LLMs and SLMs in Section 4.1 enables the modeling of the general data distributions of text $p(\text{text})$ or speech $p(\text{speech})$, the joint distribution of speech and text $p(\text{text}, \text{speech})$, or a conditional distribution $p(\text{text}|\text{speech})$ based on specific pre-training tasks. However, pre-trained models still often lack the capability to solve a large range of downstream spoken language tasks, to follow natural language instructions, or both. In the post-training phase, carefully curated datasets are used to bias SLMs toward generating desired outputs or performing tasks, typically specified using natural language instructions.

4.2.1 Task-specific training

While the eventual goal is to handle arbitrary tasks within a single model via instructions, some SLM approaches begin with a simpler post-training setting: multi-task training with task specifiers $p(\cdot|\text{speech}, \langle \text{task specifier} \rangle)$. In this approach, the pre-trained SLM is fine-tuned for a predefined set of

Task	Examples of Instructions
Speech recognition	Recognize the speech and give me the transcription. (Tang et al., 2024) Repeat after me in English. (Grattafiori et al., 2024)
Speech translation	Translate the following sentence into English. (Grattafiori et al., 2024) Recognize the speech, and translate it into English (Chu et al., 2023)
Speaker recognition	How many speakers did you hear in this audio? Who are they? (Tang et al., 2024)
Emotion recognition	Describe the emotion of the speaker. (Tang et al., 2024) Can you identify the emotion? Categorise into: sad, angry, neutral, happy (Das et al., 2024)
Question answering	What happened to this person? (Wang et al., 2023b) Generate a factual answer to preceding question (Das et al., 2024) What medicine is mentioned? Briefly introduce that medicine. (Peng et al., 2024b)

Table 2: Examples of instructions for speech-related tasks used in SLM instruction tuning.

target tasks. Qwen-Audio is an example of such an approach (Chu et al., 2023). Task-specific training can then be followed by instruction tuning or other post-training approaches, described below.

4.2.2 Instruction tuning $p(\cdot|speech, instruction)$

Instruction tuning of SLMs has been inspired by, and closely follows, successful approaches for text LLM instruction tuning. Instruction-tuning data typically consists of a speech recording, an instruction that describes the speech task, and the ground-truth output. During instruction tuning, the instructions are appended to the speech recording as inputs to the SLM, which is trained to generate the corresponding ground-truth output. Depending on the model design, the instruction can be in either text (Tang et al., 2024; Pan et al., 2024), i.e., $p(\cdot|speech, textinstruction)$, or speech format (Shu et al., 2023), i.e., $p(\cdot|speech, speechinstruction)$. In both cases, it has been found that SLMs trained on diverse instruction-tuning data can perform tasks unseen during the instruction-tuning phase (Tang et al., 2024; Das et al., 2024; Peng et al., 2024b). Table 2 shows examples of instructions for various speech tasks taken from existing instruction tuning sets.

Instruction-tuning data can be generated through various methods:

Conversion of task-specific data to instructions Task-specific data (Sec. 4.2.1) can be adapted to the instruction-tuning format by replacing task-specific tags with natural language instructions (Tang et al., 2024). LLMs can be used to rephrase those instructions to increase diversity (Arora et al., 2024).

Speech-based question answering (SQA) data Such data is typically generated using LLMs such as ChatGPT. In this process, the transcript of a speech recording is provided as input to an LLM, which is instructed to generate a question-answer pair in text format (Tang et al., 2024; Gong et al., 2024; Peng et al., 2024b). To incorporate additional context, supplementary textual descriptions about attributes such as the speaker, gender, age, emotion, and noise level may also be provided to the LLM (Yang et al., 2024b). During training, the speech recording and the corresponding LLM-generated question are provided as inputs, and the model learns to predict the LLM-generated answer.

Synthesis of text instruction-tuning data Speech instruction-tuning data can also be created by applying TTS to existing textual instruction-tuning or user-assistant conversation datasets (Peng et al., 2024b). In this approach, either a subset or the entirety of the user’s input is converted to speech. This type of data encompasses a wide variety of instruction types and response styles. The answers tend to be more descriptive than the ones in SQA and are presented in diverse textual formats, such as markdown.

Compositional instructions Instructions for individual tasks can be combined to form more complex instructions, which improves the performance on more challenging tasks. For example, an SLM can be

instructed to first perform speech recognition and then speech translation conditioned on the ASR hypothesis (Hu et al., 2024a).

4.2.3 Chat SLM training

In addition to the general post-training methods mentioned above, an important specific application setting that is gaining attention is chat SLMs, which require carefully curated training data and tailored training strategies. The development of chat SLMs has involved two main directions: (1) building a speech-aware text LM based on a text LM with chat capabilities, and (2) creating a pure speech LM or a speech+text LM that can handle speech-input-to-speech-output conversations. For (1), one approach is to use a more powerful text LLM to generate text-based conversations centered around a speech recording and use this pseudo-conversation data to fine-tune the SLM (Chu et al., 2023). For (2), a common approach is to use a TTS system to generate speech conversation data from text datasets (Zhang et al., 2023a; Défossez et al., 2024). If available, real speech conversation data can be used to provide more natural, spontaneous behaviors—such as pauses and interruptions—that occur in real conversations, though such data is often noisier and requires careful preprocessing (Défossez et al., 2024). Section 6 addresses conversational SLMs in greater detail.

4.3 Other training details

Aside from the basic strategies discussed in Sections 4.1 and 4.2, several additional training methods have proven useful for building universal SLMs, including:

Parameter-efficient fine-tuning (PEFT) Since pre-training equips LMs with strong text or speech generation capabilities, it is possible to not update the entire LM during the fine-tuning phase. Common strategies include: (1) freezing both the sequence model Seq and the speech encoder Enc^{sp} (see Figure 1) and updating only the parameters of the adaptor Adp^{sp} (Wang et al., 2023b), and (2) adding to the sequence model a set of parameter-efficient trainable adapter modules, which are trained alongside Adp^{sp} (Tang et al., 2024).

Progressive fine-tuning A key objective in training speech-aware LMs and speech+text LMs is to align the hidden representations of speech and text, in order to leverage the generation and reasoning capabilities of the pre-trained LLM. To achieve this goal, it is common to kick off post-training with content-heavy tasks, such as ASR, and gradually progress to tasks that require a more comprehensive understanding of additional information embedded in speech (e.g. emotion recognition) (Tang et al., 2024) or tasks composed of multiple sub-tasks (Hu et al., 2024b). To stabilize training, a small number of components is updated during the initial training stage, while more are updated later.

A similar strategy is adopted by conversational SLMs (Défossez et al., 2024). The intrinsic properties of human conversations present several challenges for building such SLMs. For instance, speech from different speakers often overlaps, rather than following a strictly turn-based structure. Therefore, training on real human conversations is essential. However, collecting spoken human-human conversation datasets is difficult, and publicly available datasets are often limited in size. To address these challenges, these models may be first trained on multichannel audio data before being fine-tuned on real human conversation data. See Section 6 for more details about such models.

Experience replay of pre-training data during post-training To prevent the catastrophic forgetting phenomenon (Kirkpatrick et al., 2017), reusing pre-training data during the post-training stage can help the SLM retain important capabilities learned during pre-training, such as the reasoning abilities of LLMs associated with text instruction-tuning (Peng et al., 2024b). Another approach to prevent forgetting is to use data generated by the original text LLM during post-training (Lu et al., 2024b).

Preference optimization Reinforcement learning from human feedback (RLHF) (Ouyang et al., 2022) has become a key method for aligning LLM generations with human preferences. This method was used in Qwen2-Audio (Chu et al., 2024) to improve the quality of responses generated by the speech-aware text LM. On the other hand, Align-SLM (Lin et al., 2024b) was the first to apply this method to a pure SLM and

used AI-generated feedback as a substitute for human evaluation to make the process more cost-efficient. Recently, Maimon et al. (2025) analyzed the performance of pure SLM fine-tuning using direct preference optimization as a function of fine-tuning examples.

5 Survey of representative SLMs

The previous sections have described the components, design choices, and training strategies for SLMs. In this section, we review how these choices have been instantiated in existing SLMs in the literature, categorized according to the classes in Table 1. For each category, we provide a brief review of the model design and highlight example models in the literature. Note that, aside from the examples listed, there are many SLMs with similar configurations in each category which we are unable to introduce in detail; Table 3 provides a more extensive catalogue of models and references, and Figure 6 provides a timeline to place them in historical context.

We also note that some of the multi-modal LLMs developed by industry groups, such as GPT-4o (OpenAI, 2024) and Gemini (Gemini Team et al., 2024), can also be seen as generalizations of speech+text LMs or speech-aware text LMs. These models have conversation capabilities and include text, speech, general audio, and visual inputs and outputs. However, details of their model configurations and training strategies have not been publicly released, and rigorous evaluations of these models have focused on limited speech tasks such as ASR and ST. For these reasons, it is important that future work complements these high-impact industry models with open-source, reproducible approaches.

5.1 Pure Speech LM

Pure speech LMs are trained on unlabelled speech data to model $p(\text{speech})$, and are the spoken analogue of pre-trained generative text LMs like GPT (Radford et al., 2019; Brown et al., 2020). In order to apply the next-token-prediction objectives of LM training to continuous speech data, pure speech LMs typically use discrete speech representations (Section 3.1) as the units for sequence modeling. One of the earliest approaches in this category is the generative spoken language modeling (GSLM) work of Lakhota et al. (2021). GSLM relies on a pre-trained HuBERT model $\text{Enc}^{\text{sp}}()$ to convert speech signals into phonetic tokens H^{sp} and uses an embedding layer $\text{Adp}^{\text{sp}}()$ to map token IDs to continuous vectors. The model is trained to autoregressively predict the ID of the next speech token, and the predicted tokens are converted back to audio by a unit-based vocoder $\text{Dec}^{\text{sp}}()$. GSLM’s demo⁷ shows anecdotally that this initial model can generate speech continuations with locally acceptable syntax given a speech prompt of a few seconds, but does not in general produce coherent long-form speech.

Since GSLM initiated this line of research, follow-up work has improved on it in several ways. For example, the prosody-aware generative spoken language model (pGSLM) (Kharitonov et al., 2022b) predicts prosodic information like pitch and duration jointly with phonetic tokens for better expressiveness. TWIST (Hassid et al., 2023) is similar to GSLM but the model is initialized with a pre-trained text LLM (OPT (Zhang et al., 2022)), which improves lexical and syntactic language modeling performance, as measured by the sWUGGY and sBLIMP metrics (see Section 7.1 for more details of the evaluation). Another approach (Algayres et al., 2023) uses *longer* lexical units instead of phonetic tokens, significantly reducing memory consumption. AudioLM (Borsos et al., 2023a) introduced hierarchical generation of multiple types of tokens for pure speech LMs, using a coarse-to-fine strategy that generates phonetic tokens and audio codec tokens in order (Section 3.3.1). SpeechSSM (Park et al., 2024) uses a state-space model (Gu et al., 2022) rather than a Transformer, producing more natural and semantically consistent long-form (minutes-long) speech generation.

5.2 Speech+Text LM

Speech+text LMs (Nachmani et al., 2024; Nguyen et al., 2025; Défossez et al., 2024) are SLMs that jointly model the distribution of speech and text $p(\text{text}, \text{speech})$, and can therefore understand and generate both

⁷<https://speechbot.github.io/gslm>

modalities. Such models often start with generative speech+text pre-training (Section 4.1), and can then be post-trained in a variety of ways.

A notable example of this approach is the Moshi model (Défossez et al., 2024), which is the first open-source speech-in-speech-out SLM-based dialogue system with real-time inference capability. Moshi takes time-aligned text and discrete speech representations as inputs and generates outputs in both text and speech forms. Moshi’s text processing components— $\text{Enc}^{\text{txt}}()$, $\text{Adp}^{\text{txt}}()$, and $\text{Dec}^{\text{txt}}()$, and the LM backbone $\text{Seq}()$ —are initialized from a pre-trained Helium text LLM (Défossez et al., 2024). After text pre-training, the LM is continually trained on jointly predicting the next speech and text tokens, followed by post-training in duplex mode (described in detail in Section 6) on conversation data. Moshi’s training strategy produces a model that can generate consistent speech across long (several-minute) monologues and can engage in multi-turn conversations (Défossez et al., 2024).

Speech+text LMs vary widely in their decoding methods and approaches to constructing speech and text inputs/outputs. For example, SpiRiT-LM (Nguyen et al., 2025) uses interleaved sequences of text and discrete speech tokens as both inputs and outputs, focusing solely on pre-training without doing any task-specific post-training. Compared to pure speech LMs, SpiRiT-LM has improved semantic understanding, as measured by the StoryCloze test (see Section 7.1 for more details). A precursor of this idea, SUTLM (Chou et al., 2023), modeled interleaved speech and text units of various sizes, but did not include a waveform synthesizer and was evaluated mainly on spoken language understanding tasks, with the main goal of enabling cross-modality transfer. Zeng et al. (2024b) scale up the idea of interleaved speech-text training by using a speech synthesizer to generate massive datasets of interleaved speech-text sequences. As alternatives to time-aligned or interleaved speech and text sequences, SpeechGPT (Zhang et al., 2023a) outputs a sequence of text followed by corresponding speech, while Mini-Omni (Xie & Wu, 2024a;b) and LLaMA-Omni (Fang et al., 2025) generate text and speech using separate output channels with a delay pattern in the speech token prediction (see Section 3.3.2). All three of these models (SpeechGPT, Mini-Omni, and LLaMA-Omni), unlike Moshi, focus on a traditional turn-taking structure and therefore have limited capability to model spontaneous turn-taking behavior (see Section 6 for more on conversation modeling).

5.3 Speech-aware text LM

Models in this category combine text LMs with speech encoders, and usually take both speech and text as inputs and generate responses in text form. The text input X^{txt} can include instructions asking the model to perform tasks related to the speech input X^{sp} , possibly including tasks for which the model is not specifically trained. This type of SLM is typically initialized with a text LLM (which has often already been post-trained with instruction tuning or RLHF) so as to inherit its linguistic knowledge and instruction-following capabilities. After combining the text LLM with a speech encoder and modality adapter, it is common to apply training methods that encourage better alignment between the speech and text modalities (Section 4.1.3) or to start post-training with an ASR task (Section 4.3) to help extract content information from speech inputs. The following post-training typically uses speech instruction-tuning datasets with diverse instructions (Section 4.2.2).

WavPrompt (Gao et al., 2022) represents the first example of this line of research. WavPrompt combines a wav2vec 2.0 (Baevski et al., 2020) SSL speech encoder with a GPT-2 text LM and is trained for speech classification tasks, keeping the LM backbone frozen and updating only the speech encoder. Although WavPrompt can improve over ASR+text model baselines on certain tasks, it is not evaluated on any unseen task. Since then, many models in this category have been proposed.

SALMONN (Tang et al., 2024) is another early and impactful approach. The model is built on a pre-trained Vicuna LM backbone (Zheng et al., 2023). During post-training, the model is first trained on ASR and audio captioning and later on a more diverse set of speech understanding tasks. SALMONN’s post-training keeps most of the parameters frozen—only the speech modality adapter and LoRA (Hu et al., 2022) adapters for the LM are learned—but the model is still able to generalize to unseen tasks with various natural language prompts.

Several other speech-aware text LMs were developed nearly simultaneously with SALMONN, facilitated by the release of open-source LLMs like LLaMA (Touvron et al., 2023a) and Vicuna (Zheng et al., 2023), often

post-trained and evaluated on different tasks (Gong et al., 2023b; Wang et al., 2023b; Chu et al., 2023; Chen et al., 2024d). For example, LTU-AS (Gong et al., 2023b) combines spoken content understanding with paralinguistic analysis tasks (e.g., emotion recognition). SLM (Wang et al., 2023b) generalizes to unseen tasks specified by natural language prompts even while the LM backbone and speech encoder are both frozen, and only the speech adapter is optimized during instruction-based training, but focuses only on content-heavy tasks. Qwen-Audio (Chu et al., 2023) shares similar post-training tasks to those of SALMONN, but includes learning of the speech encoder and LM backbone (sequence model) weights. During post-training, Qwen-Audio starts with a more diverse task set instead of the single ASR task, including both content and paralinguistics, followed by a chat-based supervised fine-tuning stage that improves the model’s robustness to variation in input prompts. The follow-up work, Qwen2-Audio (Chu et al., 2024) also applies direct preference optimization Rafailov et al. (2023) after supervised fine-tuning. On the other hand, the more recent DeSTA (Lu et al., 2024a) and DeSTA 2 (Lu et al., 2024b) propose a single descriptive speech-text alignment post-training task which requires the model to both recognize the spoken content and describe paralinguistic attributes. Table 4 summarizes and compares a representative set of these models’ training and evaluation choices. One key consideration is the tradeoff between preserving knowledge acquired through text pre-training and learning new speech tasks, especially those requiring non-textual information such as speech emotion. An important design choice, therefore, is which of the pre-trained parameters to freeze or update, and this issue has not yet been thoroughly explored.

In addition to the multiple post-training approaches, speech-aware text LMs also include some differences in architecture choices. For example, WavLLM (Hu et al., 2024b) combines both SSL and supervised pre-trained speech encoders, and BESTOW (Chen et al., 2024e) experiments with a sequence-to-sequence model in addition to the common decoder-only LM architecture. While most models in this category are initialized with pre-trained text LLMs, UniverSLU (Arora et al., 2024) starts from a Whisper model for conditional pre-training (see Section 4.1.2) and is then fine-tuned for various understanding tasks first by prepending new task specifiers and then by adding natural language instructions to the input.

6 Duplex speech dialogue 對話時同時說和聽

雙向語音對話 (Full-Duplex Dialogue) 的挑戰包含:

1. 語者切換時機判斷: 模型需要知道什麼時候可以講話、什麼時候該等對方講完。
2. 回饋訊號處理: 如「嗯」、「了解」這類自然語音反應, 模型需能識別與回應。
3. 被打斷的處理能力: 若模型說話時被人打斷, 需能即時做出調整。

As discussed in the previous sections, most SLMs to date assume that dialogue with a user follows a turn-by-turn structure, where a user provides an input, and the SLM generates a response. However, **natural speech dialogue is duplex**: Both speakers can simultaneously send and receive information. Compared to text-based interaction, duplex dialogue presents unique challenges. Figure 5 (a) shows an example of a full-duplex dialogue. The SLM must determine whether the user has finished their utterance before starting to speak. Backchannel signals (e.g., saying “mm-hmm,” “I see,” or “okay”) and non-verbal vocalization (e.g., laughter) are often used to facilitate smoother conversation. Additionally, the SLM may be interrupted while talking and must respond appropriately to maintain the flow of dialogue. None of these characteristics exist in text-based dialogue. Over the years, researchers have explored various methods to enable duplex dialogue in speech systems (Masumura et al., 2018; Roddy et al., 2018; Skantze, 2017; Meena et al., 2014; Ekstedt & Skantze, 2020).⁸ Rather than reviewing all prior work on spoken duplex dialogue, we focus specifically on end-to-end SLM approaches.

Dual Channel One way to achieve full-duplex dialogue is **by using dual channels** (Nguyen et al., 2023; Défossez et al., 2024; Ma et al., 2024; Meng et al., 2024), as shown in Fig. 5 (b). The SLM has two input channels: **the listening channel** (the sequence with red blocks) which continuously receives input, and **the speaking channel** (the sequence with blue blocks) where the spoken output from the SLM is directed, allowing the model to track what it has said. The model produces output at each step, including tokens representing speech or silence (blocks with a dotted outline). In this way, an SLM can listen to the user’s words while generating output simultaneously. **An early representative model for this method is the dialogue GSLM (dGSLM) (Nguyen et al., 2023); Moshi has also implemented this strategy (Défossez et al., 2024).** One challenge when using the dual-channel approach is that, instead of a typical autoregressive model, a specialized architecture is required. **dGSLM (Nguyen et al., 2023) uses a dual-tower transformer architecture,**

dGSLM、Moshi

⁸There is an extensive amount of older related literature, including on rule-based dialogue modeling; here, we only cite a few machine learning-based approaches.

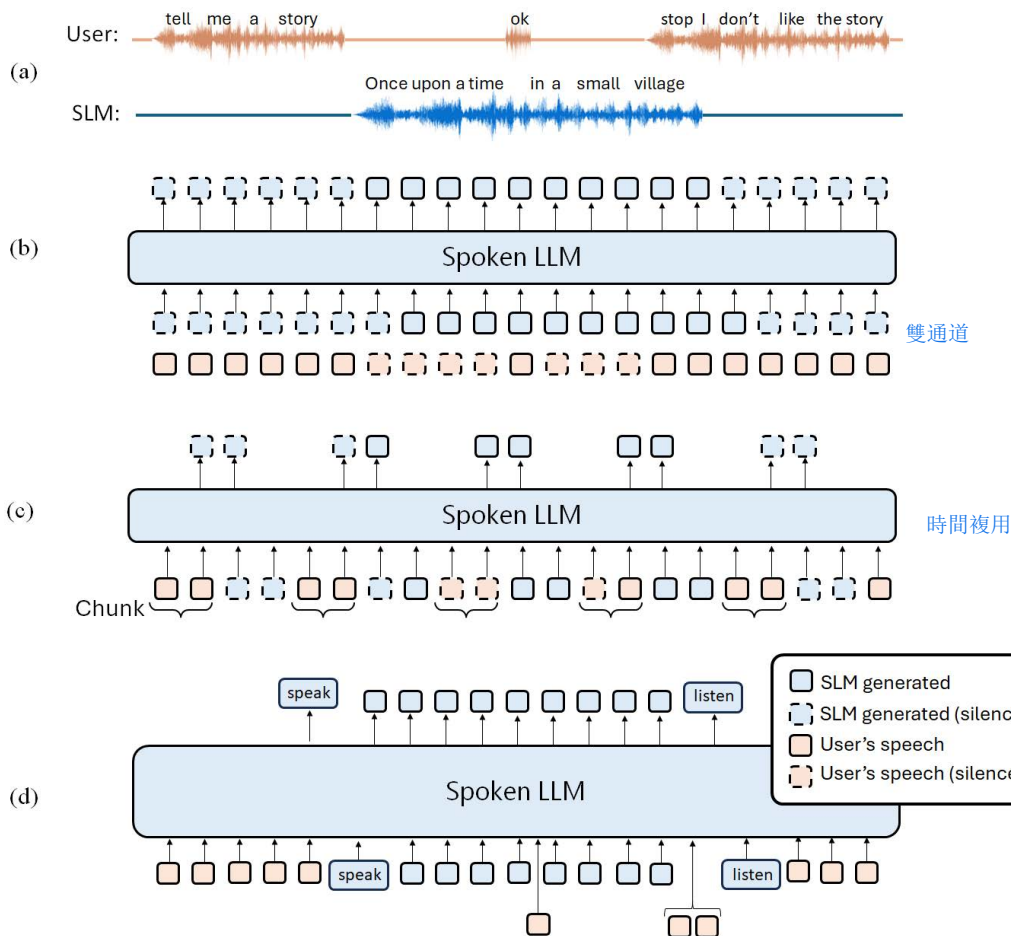


Figure 5: Full-duplex speech conversation. (a) An example of full-duplex speech conversation between a user and an SLM. (b) Dual-channel approach. (c) Time-multiplexing approach (with equal chunks). (d) Time-multiplexing approach (where the SLM controls the switching between listening and speaking modes).

where two transformers handle the two channels, using cross-attention to exchange information. Other models (Défossez et al., 2024; Ma et al., 2024; Meng et al., 2024) modify the input structure of the transformer to accommodate the dual-channel inputs.

Time Multiplexing Fig. 5 (c) and (d) illustrate the time multiplexing approach. In this approach, the SLM has only one channel, so it must switch between listening and speaking modes. During the listening mode, the SLM takes input from the user (red blocks) without generating output. During the speaking mode, the SLM generates speech representations (blue blocks) and takes its own output as input in an autoregressive manner. The strength of this approach is that the sequence model can be a typical decoder-only autoregressive model, and can therefore be initialized with a text LLM.

Synchronous LLM
、OmniFlatten,
Wang et al. 2025

The core challenge is determining how to switch between listening and speaking modes. One approach, shown in Figure 5 (c), is to alternate between processing a fixed-duration time slice from the user’s input and then switching to the speaking mode (Zhang et al., 2024c). This approach has been adopted in Synchronous LLM (Veluri et al., 2024) and OmniFlatten (Zhang et al., 2025).

Another approach allows the model to determine when to switch modes (Wang et al., 2024b; Xu et al., 2024; Wang et al., 2024c), as shown in Figure 5 (d). In listening mode, while processing the user’s input, at each time step the model predicts whether to initiate a response. This decision is signalled by the prediction of a special token ([speak], as shown in the figure). The model processes its generated output auto-regressively

優點是：可以使用一般的 decoder-only 自回歸模型，因此可以直接用一個文字語言模型（LLM）初始化

until it produces the [listen] token, which triggers a switch back to listening mode. Users can continue providing input during speaking mode, dynamically influencing the model’s responses. In particular, user input influences the [listen] token generation, enabling users to interrupt the model’s output when desired. This interaction is implemented by interleaving the model’s responses with user input, using some special tokens or embeddings to distinguish between the two sources.⁹

7 Benchmarking and Evaluating SLMs

由於語音比文字複雜，
不容易用單一標準來比較各種 SLMs，
評估時更需要從多角度、多階段訓練層面來審視它們

Like text language models, SLMs may possess a broad array of capabilities. In addition (and in contrast to text LMs), they also involve a variety of design choices that differ significantly across models. These properties can make it difficult to compare SLMs, and more important to assess them from multiple angles and training stages. Below, we outline commonly used evaluation methods and benchmarks, categorized into three primary subgroups: (1) **likelihood-based evaluation metrics**; (2) **generative metrics**; and (3) **trustworthiness evaluations**. Traditional task-specific speech evaluation methods (e.g., ASR, TTS, ST) are also frequently used to assess SLM performance. However, we exclude them from this survey, as they are well-established and extensively covered in the existing literature.

7.1 Likelihood-Based Evaluation

As described above, the general SLM architecture comprises three main components: (1) speech (and optionally text) encoding; (2) sequence modeling of speech, text, or both; and (3) speech and/or text decoding. Likelihood-based evaluation metrics (Dunbar et al., 2021) are designed to assess the sequential speech modeling capabilities of SLMs.¹⁰ Such evaluation metrics are therefore more suitable for the pre-training stage of pure speech LMs and joint speech+text LMs (Lakhotia et al., 2021; Hassid et al., 2023; Borsos et al., 2023a; Défossez et al., 2024).

It is crucial to emphasize that speech encoding (i.e., speech tokenization), an active area of research (Messica & Adi, 2024; Turetzky & Adi, 2024; Zhang et al., 2024b), plays a significant role in shaping speech modeling capabilities and therefore cannot be disregarded. For instance, utilizing tokens generated by a self-supervised learning model such as HuBERT significantly outperforms the quantization of log mel spectrograms (Lakhotia et al., 2021). Likewise, employing speech tokenizers with higher compression rates (i.e., sampled at lower frequency in the latent space) yields better performance in speech modeling (Borsos et al., 2023a; Hassid et al., 2023). While this subsection discusses speech modeling tasks with a focus on the LM, we believe that in the context of SLMs, the research community should evaluate both components collectively rather than treating them as separate entities.

In likelihood-based evaluations, the pre-trained SLM is typically provided with two input sequences: one representing a natural speech sequence (positive example) and the other an unnatural sequence relative to a specific property (negative example). We then calculate the percentage of instances where the SLM assigns a higher likelihood to the positive example than to the negative one. This type of evaluation has also been popular in the NLP community until very recently (Zellers et al., 2019; Mostafazadeh et al., 2016; Touvron et al., 2023b; Gemma Team et al., 2024; LCM team et al., 2024). Due to significant model improvements this type of NLP benchmark has saturated and significantly more complex ones have emerged (Jimenez et al., 2023; Srivastava et al., 2023).

Metrics in this category are designed to evaluate a model’s capacity to represent various speech characteristics. For example, sWUGGY (Dunbar et al., 2021) evaluates the lexical capabilities of models by presenting them with pairs of utterances, one consisting of a real word and the other a similar non-word (e.g., ‘brick’ vs. ‘blick’), and measuring the model’s ability to assign a higher probability to the real word (where nonwords were obtained from the text WUGGY task (Keuleers & Brysbaert, 2010)). Similarly, sBLIMP (Dunbar et al., 2021) (which was adapted from the text-based BLIMP (Warstadt et al., 2020)) evaluates the ability of the

⁹Theoretically, there is always some input, even when no one is speaking. Here, we assume that a voice activity detection (VAD) system is in place, so only the tokens representing actual user speech are sent to the SLM during speaking mode.

¹⁰Such evaluation metrics can be considered as an analogue to perplexity (PPL) for text LLMs. Since there is no standard method for speech tokenization, the speech “vocabulary” may change between pure SLMs. Therefore, PPL cannot be directly applied. Instead, likelihood-based metrics introduce an alternative approach of comparing sequence likelihoods.

network to model syntactic properties. The network is presented with a matched pair of grammatically correct and incorrect sentences. Hassid et al. (2023) expanded this method to evaluate semantic understanding of spoken text. They generated two spoken adaptations of the StoryCloze text paragraph story completion task (Mostafazadeh et al., 2016). In the first adaptation, they adhered to the original StoryCloze setup, while in the second, they randomly sampled negative examples, resulting in a simplified version primarily requiring an understanding of the general topic of the paragraph. ProsAudit (de Seyssel et al., 2023) constructs negative examples by inserting pauses in unnatural positions in utterances to evaluate SLMs’ sensitivity to speech prosody. Lastly, Maimon et al. (2024) proposed the SALMon evaluation suite, which, unlike the previously mentioned benchmarks, focuses on a variety of non-lexical acoustic and prosodic elements (e.g., speaker identity, speaker sentiment, background noise and room acoustics) at two levels of complexity: One measures the *consistency* across time of a given acoustic element, whereas the other measures *alignment* between the acoustic details and the semantic, spoken content.

7.2 Generative Metrics

Unlike likelihood-based metrics, which assess a model’s capacity to assign higher likelihoods to in-domain samples than out-of-domain ones, generative metrics focus on evaluating the model’s ability to produce meaningful continuations or responses based on a given prompt or instruction. Such evaluation methods can be suitable for all SLM types (pure speech LMs, speech+text LMs, or speech-aware text LMs).

Intrinsic quality metrics For SLMs that generate spoken output, the quality of the generated response can be evaluated along several axes: (1) speech quality, using subjective tests or objective metrics such as MOSNet (Lo et al., 2019); (2) speaker, acoustic, and/or sentiment consistency using human judges or pre-trained classifiers (Nguyen et al., 2025); and (3) quality of the spoken content, evaluated by either comparing to a target response or transcribing the speech and using a judge or text LLM to score the content.

For pure SLMs, a common evaluation metric is the *generative perplexity*, in which we transcribe the generated speech and measure its perplexity (PPL) using a pre-trained text LLM (Lakhotia et al., 2021). However, a known problem with this metric is that text LLMs tend to assign high probability to repeated content (Holtzman et al., 2020). To mitigate this issue, Lakhotia et al. (2021) propose the VERT metric, which balances between quality and diversity of the generated response. VERT uses a linear combination of the text PPL and the auto-BLEU metric of within-sentence diversity (Lakhotia et al., 2021).

Task-based evaluation For joint speech+text LMs and speech-aware text LMs, a common approach is to evaluate the model’s ability to perform question answering and follow instructions. Nachmani et al. (2024) proposed two spoken question answering benchmarks: LLaMA-Questions, derived from a text LLM, and a synthesized version of the WebQuestions (Berant et al., 2013) textual benchmark. Défossez et al. (2024) built upon this approach and developed a synthesized version of TriviaQA (Joshi et al., 2017). To evaluate SLM responses on these question-answering benchmarks, the generated speech is transcribed and the accuracy is measured against the target answer. Chen et al. (2024c) proposed VoiceBench, an extended question-answering dataset consisting of (1) open-ended QA; (2) reference-based QA; (3) multiple-choice QA; and (4) instruction-following evaluations. VoiceBench also evaluates the model’s ability to handle background noise and speaker variability. Fang et al. (2025) proposed to evaluate the instruction-following capabilities of SLMs using *LLM-as-a-judge* methods. In this approach, an external text LLM rates the quality of the responses considering both content and style.

Recently, there have been several efforts to develop evaluations of speech-aware text LMs on a broader range of tasks and instruction data. In these evaluations, each task consists of text instructions, speech utterances, and text labels. Huang et al. (2024) introduced the Dynamic-SUPERB evaluation suite, designed as a dynamic benchmark allowing for community contributions of tasks, analogously to BIG-bench (Srivastava et al., 2023) for text LMs. The initial phase of Dynamic-SUPERB focused on classification tasks related to content, speaker characteristics, semantics, degradation, paralinguistics, and non-speech audio information. Phase 2 of Dynamic-SUPERB (Huang et al., 2025) significantly expanded the task set to include regression and sequence generation tasks, making it the largest benchmark for speech and audio evaluation. Yang et al. (2024b) proposed AIR-Bench, which includes both classification and open-ended chat-style questions. Both

Dynamic-SUPERB Phase-2 and AIR-Bench use LLM-as-a-judge evaluation. Sakshi et al. (2025) introduced MMAU, a benchmark for audio understanding and reasoning using natural language. MMAU includes both information extraction tasks (e.g., “Which word appears first?”) and reasoning tasks (e.g., “What are the roles of the first and second speakers in the conversation?”), formulated as multiple-choice questions. Finally, AudioBench (Wang et al., 2024a) includes eight speech and audio tasks based on 26 datasets, focusing on speech understanding (e.g., ASR, spoken QA, speech instruction), audio scene comprehension (e.g., audio captioning, audio scene question answering), and voice analysis (e.g., accent recognition, gender recognition, emotion recognition).

While the benchmarks above include a large variety of tasks, some recent benchmarks focus specifically on speech characteristics that exploit the new abilities of SLMs. For example, StyleTalk (Lin et al., 2024a) assesses SLMs’ ability to generate text responses that align with specific speaking styles, such as intonation and emotion. The evaluation compares the generated output with a target response using standard text generation metrics (BLEU, ROUGE, and others). Similarly, the E-chat200 dataset (Xue et al., 2024) also assesses a model’s ability to generate responses that align with the speaker’s emotion. The dataset was synthetically created using an expressive TTS, while the text-based questions and responses were generated by an external LLM. SD-Eval (Ao et al., 2024) goes beyond emotion alignment to include alignment with paralinguistic features (e.g., accent, age, prosody, timbre, volume) as well as environmental context.

7.3 Evaluating Trustworthiness

The benchmarks discussed thus far focus on the *performance* of SLMs. Next, we consider evaluations that measure their *trustworthiness*. Since SLMs generate word sequences (whether directly as text or within the spoken output), all trustworthiness considerations related to text-based LLMs also apply to SLMs. However, in addition to this content information, non-content information in the generated speech creates additional aspects of trustworthiness that need to be considered.

7.3.1 Hallucination

Text-based LLMs often suffer from hallucinations, or outputs that are inconsistent with the given context or other knowledge considered to be fact. In addition to such content hallucinations, SLMs may also generate audio hallucinations. For example, consider an audio clip without a dog barking. If you ask an SLM to describe it, it would most likely not mention anything related to a dog barking. However, if you directly ask the model, “Is there a dog barking?” it might answer, “Yes,” even though it can provide accurate audio captions when explicitly prompted to do so (Kuan & Lee, 2024).

Kuan & Lee (2024) study whether SLMs can accurately perform three types of tasks without hallucination: (1) identify the presence of an object (e.g., “Is there a dog barking?”), (2) determine the temporal order of sound events (e.g., “Is the dog barking before someone laughs?”), and (3) recognize the source of a sound event (e.g., “Who is laughing, a man or a woman?”). The main finding is that all of the evaluated SLMs hallucinate more than a simple cascade model that combines audio captioning with text LLMs. This issue may arise because, while SLMs can generate accurate audio captions, they struggle to understand specific questions Kuan et al. (2024). The authors also proposed a method that mitigates hallucination by prompting the model to produce output in multiple steps (similarly to chain-of-thought prompting (Wei et al., 2022b)), by first describing the audio and then responding to the instruction based on that description.

7.3.2 Toxicity

Prevention of harmful, offensive, or inappropriate output is a crucial issue for text LLMs, and the same challenge also applies to SLMs. Both Meta’s SpiRit-LM (Nguyen et al., 2025) and OpenAI’s GPT-4o voice mode (OpenAI, 2024) evaluate the toxicity of their models’ responses. These evaluations typically involve using specific prompts to elicit speech responses from the SLMs and assessing the toxicity of the text transcriptions of those responses. These evaluations therefore consider only verbal toxicity, i.e. toxicity within the word sequence. The evaluation of non-verbal toxicity (e.g., toxic sarcasm) has yet to be widely studied.

7.3.3 Bias

Similarly to toxicity, all bias measures used for text LLMs can also be applied to SLMs. For SLMs that produce spoken output, text-based methods can still be used to analyze their transcriptions (Lin et al., 2024d). In addition, because speech conveys information beyond the textual content, there may be factors outside the text that indicate bias. Lin et al. (2024c) investigate such biases by providing an SLM with input utterances that share identical content but differ in speaker characteristics (such as gender and age) to assess whether these traits influence its responses. The results suggest that the tested SLMs exhibit minimal bias. However, this may be because current models lack sufficient sophistication to recognize differences in speaker characteristics, and therefore do not respond differently based on them. OpenAI report in a blog post (OpenAI, 2024) on their efforts to ensure consistency across user voices by post-training GPT-4o with a diverse set of voices and on a bias evaluation, which found no significant variation in model behavior across voices. While such findings are encouraging, regular and thorough bias benchmarking is needed to track the evolution of SLM bias as this field matures.

7.3.4 Deepfakes

Current SLMs can mimic a variety of voices to a level indistinguishable from human speech (Chen et al., 2025b). Unfortunately, malicious actors may exploit this technology, leading to misuse and security concerns, such as deepfake attacks (for example, creating fake news using the voice of a public figure). To address this issue, (Wu et al., 2024e; Du et al., 2025) introduced CodecFake, a deepfake audio dataset, and found that detection models trained on CodecFake can counter fake audio generated by existing SLMs.

8 Challenges and future work

This survey has covered some of the key milestones that have been achieved in SLM research: the first pure speech language models that could generate convincing-sounding stretches of English speech (GSLM (Lakhotia et al., 2021)); the first generation of models that could perform a variety of tasks with reasonable accuracy given natural language instructions (Gao et al., 2022; Gong et al., 2023b; Wang et al., 2023b); models that can converse in “full duplex” mode (Défossez et al., 2024); and benchmarks tailored for evaluating the modeling and instruction-following capabilities of SLMs (Dunbar et al., 2021; Huang et al., 2025).

While research in this area has produced many SLMs and exciting outcomes, it is safe to say that we are not yet close to the goal of truly universal speech processing systems. This section highlights several categories of challenges and open questions, which suggest directions for future research.

Model architecture The optimal representation of speech within SLMs remains unclear. Speech representations in SLMs include both discrete and continuous varieties, derived from a wide range of encoders. This design choice can also influence other architectural choices in an SLM, for example depending on the information rate of the encoder and whether it encodes more phonetic or other types of information.

Another open question is determining the best method to combine speech and text, which applies to all aspects of SLM modeling and training. We have described various choices of modality adapters and approaches for interleaving speech and text. These have not been thoroughly compared, so the effect of each modeling choice is still unclear.

A final architectural challenge is that current SLMs are large and slow, making them impractical for real-time and on-device settings. To some extent this is because various compression algorithms (e.g., (Lai et al., 2021; Peng et al., 2023a; Ding et al., 2024)) and alternative architectures (e.g., Park et al. (2024)) have not been widely applied to SLMs. However, there is also an inherent efficiency challenge that arises when combining multiple pre-trained components, sometimes with different architectures and frame rates.

Training There is a lack of public high-quality training data, particularly for instruction tuning and chat-based training. Zeng et al. (2024b) showed that scaling synthetic data generation enhances the performance of pure speech LMs; similar research efforts could be directed toward instruction tuning and chat-based spoken data. Additionally, SLMs are trained on diverse datasets (including some proprietary datasets),

making it difficult to analyze whether performance differences are caused by architecture design choices or training data. Thorough ablation studies focusing on the various model design choices (Section 3) and training strategies (Section 4) are also essential and still lacking. Finally, scaling studies for SLMs are needed in order to better understand how SLMs scale with model and data. Such studies would hopefully produce new scaling laws and help to speed up the development cycle Cuervo & Marxer (2024) conducted the first such investigation for pure SLMs. More recently, Maimon et al. (2025) presented a thorough empirical analysis of the optimization process for pure speech LMs, showing how one can train a high-quality pure SLM in 24 hours using a single GPU. It remains uncertain whether scaling findings apply also to speech+text LMs or speech-aware text LMs.

Evaluation Thus far, different SLMs have typically been evaluated on different datasets and tasks. Recent efforts to collect large numbers of tasks and standardize benchmarking (Huang et al., 2024; Yang et al., 2024b; Huang et al., 2025) are promising, but as of this writing they are not yet widely adopted for evaluating new models. Existing benchmarks also do not cover the full range of spoken language tasks. The largest current benchmarking effort, Dynamic-SUPERB Phase-2 (Huang et al., 2025), includes 180 tasks, which is similar to the size of the BIG-Bench suite of text LLM tasks (Srivastava et al., 2023). However, the range of spoken language tasks is arguably much larger than that of text tasks, since speech tasks include the vast majority of text tasks (the ones that don’t relate explicitly to the textual form, such as transliteration) and in addition include a variety of speech-specific tasks related to speaker properties, accents, or prosody-specific content. In addition, there is a lack of standardized benchmarking for speech *generation* tasks.

Open research Very few SLMs are fully open-source—including code, model checkpoints, and training data—which makes a comprehensive comparison between approaches virtually impossible. There has been progress in this direction, with some models having at least publicly available weights (see Fig. 6), and some support in open-source toolkits (Tian et al., 2025). Many of the items on the wish list above, such as controlled comparisons between multiple design decisions, will be infeasible to accomplish until more fully open models are released.

Improving inclusiveness and safety SLM research has, thus far, understandably focused on high-resource languages and settings. As SLMs become more performant and enter commercial products in daily use, it will be critical to make them accessible to as broad a range of users as possible, including a variety of languages, dialects, and speech-related medical conditions. Research in this area will likely follow the arc of text LLM research, but again, there are speech-specific challenges that arise from speaker variation. Similarly, the area of safety and trustworthiness has only begun to be explored, and will require speech-specific solutions to speech-specific challenges like deepfakes and speaker type-related bias.

References

- Andrea Agostinelli, Timo I. Denk, Zalán Borsos, Jesse Engel, Mauro Verzetti, Antoine Caillon, Qingqing Huang, Aren Jansen, Adam Roberts, Marco Tagliasacchi, Matt Sharifi, Neil Zeghidour, and Christian Frank. MusicLM: Generating music from text. *arXiv preprint arXiv:2301.11325*, 2023.
- Robin Algayres, Adel Nabli, Benoît Sagot, and Emmanuel Dupoux. Speech sequence embeddings using nearest neighbors contrastive learning. In *Proc. Interspeech*, 2022.
- Robin Algayres, Yossi Adi, Tu Nguyen, Jade Copet, Gabriel Synnaeve, Benoît Sagot, and Emmanuel Dupoux. Generative spoken language model based on continuous word-sized audio tokens. In *Proc. EMNLP*, 2023.
- Keyu An, Qian Chen, Chong Deng, Zhihao Du, Changfeng Gao, Zhifu Gao, Yue Gu, Ting He, Hangrui Hu, Kai Hu, et al. Funaudiollm: Voice understanding and generation foundation models for natural interaction between humans and llms. *arXiv preprint arXiv:2407.04051*, 2024.
- Junyi Ao, Yuancheng Wang, Xiaohai Tian, Dekun Chen, Jun Zhang, Lu Lu, Yuxuan Wang, Haizhou Li, and Zhizheng Wu. SD-Eval: A benchmark dataset for spoken dialogue understanding beyond words. In *Proc. NeurIPS*, 2024.

-
- Siddhant Arora, Hayato Futami, Jee-weon Jung, Yifan Peng, Roshan Sharma, Yosuke Kashiwagi, Emiru Tsunoo, Karen Livescu, and Shinji Watanabe. UniverSLU: Universal spoken language understanding for diverse tasks with natural language instructions. In *Proc. NAACL*, 2024.
- Alexei Baevski, Yuhao Zhou, Abdelrahman Mohamed, and Michael Auli. wav2vec 2.0: A framework for self-supervised learning of speech representations. In *Proc. NeurIPS*, 2020.
- Jinze Bai, Shuai Bai, Yunfei Chu, Zeyu Cui, Kai Dang, Xiaodong Deng, Yang Fan, Wenbin Ge, Yu Han, Fei Huang, Binyuan Hui, Luo Ji, Mei Li, Junyang Lin, Runji Lin, Dayiheng Liu, Gao Liu, Chengqiang Lu, Keming Lu, Jianxin Ma, Rui Men, Xingzhang Ren, Xuancheng Ren, Chuanqi Tan, Sinan Tan, Jianhong Tu, Peng Wang, Shijie Wang, Wei Wang, Shengguang Wu, Benfeng Xu, Jin Xu, An Yang, Hao Yang, Jian Yang, Shusheng Yang, Yang Yao, Bowen Yu, Hongyi Yuan, Zheng Yuan, Jianwei Zhang, Xingxuan Zhang, Yichang Zhang, Zhenru Zhang, Chang Zhou, Jingren Zhou, Xiaohuan Zhou, and Tianhang Zhu. Qwen technical report. *arXiv preprint arXiv:2309.16609*, 2023.
- Jonathan Berant, Andrew Chou, Roy Frostig, and Percy Liang. Semantic parsing on Freebase from question-answer pairs. In *Proc. EMNLP*, 2013.
- Zalán Borsos, Raphaël Marinier, Damien Vincent, Eugene Kharitonov, Olivier Pietquin, Matt Sharifi, Dominik Roblek, Olivier Teboul, David Grangier, Marco Tagliasacchi, and Neil Zeghidour. AudioLM: A language modeling approach to audio generation. *IEEE/ACM Trans. Audio, Speech, Lang. Process.*, 31: 2523–2533, 2023a.
- Zalán Borsos, Matt Sharifi, Damien Vincent, Eugene Kharitonov, Neil Zeghidour, and Marco Tagliasacchi. SoundStorm: Efficient parallel audio generation. *arXiv preprint arXiv:2305.09636*, 2023b.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Nee-lakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. In *Proc. NeurIPS*, 2020.
- William Chan, Navdeep Jaitly, Quoc V. Le, and Oriol Vinyals. Listen, attend and spell. In *Proc. ICASSP*, 2016.
- Huiwen Chang, Han Zhang, Lu Jiang, Ce Liu, and William T. Freeman. MaskGIT: Masked generative image transformer. In *Proc. CVPR*, 2022.
- Kai-Wei Chang, Haibin Wu, Yu-Kai Wang, Yuan-Kuei Wu, Hua Shen, Wei-Cheng Tseng, Iu-Thing Kang, Shang-Wen Li, and Hung-yi Lee. SpeechPrompt: Prompting speech language models for speech processing tasks. *IEEE/ACM Trans. Audio, Speech, Lang. Process.*, 32:3730–3744, 2024.
- Feilong Chen, Minglun Han, Haozhi Zhao, Qingyang Zhang, Jing Shi, Shuang Xu, and Bo Xu. X-LLM: Bootstrapping advanced large language models by treating multi-modalities as foreign languages. *arXiv preprint arXiv:2305.04160*, 2023.
- Liang Chen, Zekun Wang, Shuhuai Ren, Lei Li, Haozhe Zhao, Yunshui Li, Zefan Cai, Hongcheng Guo, Lei Zhang, Yizhe Xiong, Yichi Zhang, Ruoyu Wu, Qingxiu Dong, Ge Zhang, Jian Yang, Lingwei Meng, Shujie Hu, Yulong Chen, Junyang Lin, Shuai Bai, Andreas Vlachos, Xu Tan, Minjia Zhang, Wen Xiao, Aaron Yee, Tianyu Liu, and Baobao Chang. Next token prediction towards multimodal intelligence: A comprehensive survey. *arXiv preprint arXiv:2412.18619*, 2024a.
- Qian Chen, Yafeng Chen, Yanni Chen, Mengzhe Chen, Yingda Chen, Chong Deng, Zhihao Du, Ruize Gao, Changfeng Gao, Zhifu Gao, et al. Minmo: A multimodal large language model for seamless voice interaction. *arXiv preprint arXiv:2501.06282*, 2025a.

-
- Sanyuan Chen, Chengyi Wang, Zhengyang Chen, Yu Wu, Shujie Liu, Zhuo Chen, Jinyu Li, Naoyuki Kanda, Takuya Yoshioka, Xiong Xiao, Jian Wu, Long Zhou, Shuo Ren, Yanmin Qian, Yao Qian, Jian Wu, Michael Zeng, Xiangzhan Yu, and Furu Wei. WavLM: Large-scale self-supervised pre-training for full stack speech processing. *IEEE J. Sel. Top. Signal Process.*, 16(6):1505–1518, 2022.
- Sanyuan Chen, Chengyi Wang, Yu Wu, Ziqiang Zhang, Long Zhou, Shujie Liu, Zhuo Chen, Yanqing Liu, Huaming Wang, Jinyu Li, Lei He, Sheng Zhao, and Furu Wei. Neural codec language models are zero-shot text to speech synthesizers. *IEEE Trans. Audio, Speech, Lang. Process.*, 33:705–718, 2025b.
- Wenxi Chen, Ziyang Ma, Ruiqi Yan, Yuzhe Liang, Xiquan Li, Ruiyang Xu, Zhikang Niu, Yanqiao Zhu, Yifan Yang, Zhanxun Liu, et al. Slam-omni: Timbre-controllable voice interaction system with single-stage training. *arXiv preprint arXiv:2412.15649*, 2024b.
- Yiming Chen, Xianghu Yue, Chen Zhang, Xiaoxue Gao, Robby T. Tan, and Haizhou Li. VoiceBench: Benchmarking LLM-based voice assistants. *arXiv preprint arXiv:2410.17196*, 2024c.
- Zhehuai Chen, He Huang, Andrei Andrusenko, Oleksii Hrinchuk, Krishna C Puvvada, Jason Li, Subhankar Ghosh, Jagadeesh Balam, and Boris Ginsburg. SALM: Speech-augmented language model with in-context learning for speech recognition and translation. In *Proc. ICASSP*, 2024d.
- Zhehuai Chen, He Huang, Oleksii Hrinchuk, Krishna C. Puvvada, Nithin Rao Koluguri, Piotr Żelasko, Jagadeesh Balam, and Boris Ginsburg. Bestow: Efficient and streamable speech language model with the best of two worlds in GPT and T5. In *Proc. SLT*, 2024e.
- Ju-Chieh Chou, Chung-Ming Chien, Wei-Ning Hsu, Karen Livescu, Arun Babu, Alexis Conneau, Alexei Baevski, and Michael Auli. Toward joint language modeling for speech units and text. In *Findings of EMNLP*, 2023.
- Yunfei Chu, Jin Xu, Xiaohuan Zhou, Qian Yang, Shiliang Zhang, Zhijie Yan, Chang Zhou, and Jingren Zhou. Qwen-Audio: Advancing universal audio understanding via unified large-scale audio-language models. *arXiv preprint arXiv:2311.07919*, 2023.
- Yunfei Chu, Jin Xu, Qian Yang, Haojie Wei, Xipin Wei, Zhifang Guo, Yichong Leng, Yuanjun Lv, Jinzheng He, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen2-Audio technical report. *arXiv preprint arXiv:2407.10759*, 2024.
- Jade Copet, Felix Kreuk, Itai Gat, Tal Remez, David Kant, Gabriel Synnaeve, Yossi Adi, and Alexandre Défossez. Simple and controllable music generation. *Proc. NeurIPS*, 2023.
- Santiago Cuervo and Ricard Marxer. Scaling properties of speech language models. In *Proc. EMNLP*, 2024.
- Wenqian Cui, Dianzhi Yu, Xiaoqi Jiao, Ziqiao Meng, Guangyan Zhang, Qichao Wang, Yiwen Guo, and Irwin King. Recent advances in speech language models: A survey. *arXiv preprint arXiv:2410.03751*, 2024.
- Nilaksh Das, Saket Dingliwal, Srikanth Ronanki, Rohit Paturi, David Huang, Prashant Mathur, Jie Yuan, Dhanush Bekal, Xing Niu, Sai Muralidhar Jayanthi, Xilai Li, Karel Mundnich, Monica Sunkara, Sundararajan Srinivasan, Kyu J. Han, and Katrin Kirchhoff. SpeechVerse: A large-scale generalizable audio language model. *arXiv preprint arXiv:2405.08295*, 2024.
- Maureen de Seyssel, Marvin Lavechin, Hadrien Titeux, Arthur Thomas, Gwendal Virlet, Andrea Santos Revilla, Guillaume Wisniewski, Bogdan Ludusan, and Emmanuel Dupoux. Prosaudit, a prosodic benchmark for self-supervised speech models. In *Proc. Interspeech*, 2023.
- Alexandre Défossez, Jade Copet, Gabriel Synnaeve, and Yossi Adi. High fidelity neural audio compression. *Trans. MLR*, 2023.
- Alexandre Défossez, Laurent Mazaré, Manu Orsini, Amélie Royer, Patrick Pérez, Hervé Jégou, Edouard Grave, and Neil Zeghidour. Moshi: A speech-text foundation model for real-time dialogue. *arXiv preprint arXiv:2410.00037*, 2024.

-
- Keqi Deng, Guangzhi Sun, and Philip C. Woodland. Wav2Prompt: End-to-end speech prompt generation and tuning for LLM in zero and few-shot learning. *arXiv preprint arXiv:2406.00522*, 2024.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proc. NAACL*, 2019.
- Shaojin Ding, David Qiu, David Rim, Yanzhang He, Oleg Rybakov, Bo Li, Rohit Prabhavalkar, Weiran Wang, Tara N. Sainath, Zhonglin Han, Jian Li, Amir Yazdanbakhsh, and Shivani Agrawal. USM-Lite: Quantization and sparsity aware fine-tuning for speech recognition with universal speech models. In *Proc. ICASSP*, 2024.
- Jiawei Du, Xuanjun Chen, Haibin Wu, Lin Zhang, I-Ming Lin, I-Hsiang Chiu, Wenze Ren, Yuan Tseng, Yu Tsao, Jyh-Shing Roger Jang, and Hung-yi Lee. CodecFake-Omni: A large-scale codec-based deepfake speech dataset. *arXiv preprint arXiv:2501.08238*, 2025.
- Zhihao Du, Qian Chen, Shiliang Zhang, Kai Hu, Heng Lu, Yexin Yang, Hangrui Hu, Siqi Zheng, Yue Gu, Ziyang Ma, Zhifu Gao, and Zhijie Yan. CosyVoice: A scalable multilingual zero-shot text-to-speech synthesizer based on supervised semantic tokens. *arXiv preprint arXiv:2407.05407*, 2024.
- Ewan Dunbar, Mathieu Bernard, Nicolas Hamilakis, Tu Anh Nguyen, Maureen De Seyssel, Patricia Rozé, Morgane Rivière, Eugene Kharitonov, and Emmanuel Dupoux. The zero resource speech challenge 2021: Spoken language modelling. In *Proc. Interspeech*, 2021.
- Erik Ekstedt and Gabriel Skantze. TurnGPT: A transformer-based language model for predicting turn-taking in spoken dialog. In *Findings of ACL*, 2020.
- Ruchao Fan, Bo Ren, Yuxuan Hu, Rui Zhao, Shujie Liu, and Jinyu Li. AlignFormer: Modality matching can achieve better zero-shot instruction-following speech-LLM. *arXiv preprint arXiv:2412.01145*, 2024.
- Qingkai Fang, Shoutao Guo, Yan Zhou, Zhengrui Ma, Shaolei Zhang, and Yang Feng. LLaMA-Omni: Seamless speech interaction with large language models. In *Proc. ICLR*, 2025.
- Yassir Fathullah, Chunyang Wu, Egor Lakomkin, Ke Li, Junteng Jia, Yuan Shangguan, Jay Mahadeokar, Ozlem Kalinli, Christian Fuegen, and Mike Seltzer. AudioChatLlama: Towards general-purpose speech abilities for LLMs. In *Proc. NAACL*, 2024.
- Philip Gage. A new algorithm for data compression. *C Users J.*, 12(2):23–38, 1994.
- Marco Gaido, Mauro Cettolo, Matteo Negri, and Marco Turchi. CTC-based compression for direct speech translation. In *Proc. EACL*, 2021.
- Heting Gao, Junrui Ni, Kaizhi Qian, Yang Zhang, Shiyu Chang, and Mark Hasegawa-Johnson. WavPrompt: Towards few-shot spoken language understanding with frozen language models. In *Proc. Interspeech*, 2022.
- Gemini Team et al. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context. *arXiv preprint arXiv:2403.05530*, 2024.
- Gemma Team et al. Gemma: Open models based on gemini research and technology. *arXiv preprint arXiv:2403.08295*, 2024.
- Yuan Gong, Sameer Khurana, Leonid Karlinsky, and James Glass. Whisper-AT: Noise-robust automatic speech recognizers are also strong general audio event taggers. In *Proc. Interspeech*, 2023a.
- Yuan Gong, Alexander H. Liu, Hongyin Luo, Leonid Karlinsky, and James Glass. Joint audio and speech understanding. In *Proc. ASRU*, 2023b.
- Yuan Gong, Hongyin Luo, Alexander H. Liu, Leonid Karlinsky, and James R. Glass. Listen, think, and understand. In *Proc. ICLR*, 2024.
- Aaron Grattafiori et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

-
- Alex Graves, Santiago Fernández, Faustino Gomez, and Jürgen Schmidhuber. Connectionist temporal classification: Labelling unsegmented sequence data with recurrent neural networks. In *Proc. ICML*, 2006.
- R. Gray. Vector quantization. *IEEE ASSP Magazine*, 1(2):4–29, 1984.
- Albert Gu, Karan Goel, and Christopher Ré. Efficiently modeling long sequences with structured state spaces. In *Proc. ICLR*, 2022.
- Yiwei Guo, Zhihan Li, Hankun Wang, Bohan Li, Chongtian Shao, Hanglei Zhang, Chenpeng Du, Xie Chen, Shujie Liu, and Kai Yu. Recent advances in discrete speech tokens: A review. *arXiv preprint arXiv:2502.06490*, 2025.
- Michael Hassid, Tal Remez, Tu Anh Nguyen, Itai Gat, Alexis Conneau, Felix Kreuk, Jade Copet, Alexandre Défossez, Gabriel Synnaeve, Emmanuel Dupoux, Roy Schwartz, and Yossi Adi. Textually pretrained speech language models. In *Proc. NeurIPS*, 2023.
- William Held, Ella Li, Michael J. Ryan, Weiyan Shi, Yanzhe Zhang, and Diyi Yang. Distilling an end-to-end voice assistant without instruction training data. *arXiv preprint arXiv:2410.02678*, 2024.
- François Hernandez, Vincent Nguyen, Sahar Ghannay, Natalia Tomashenko, and Yannick Esteve. Ted-lium 3: Twice as much data and corpus repartition for experiments on speaker adaptation. In *Speech and Computer: 20th International Conference, SPECOM 2018, Leipzig, Germany, September 18–22, 2018, Proceedings 20*, pp. 198–208. Springer, 2018.
- Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration. In *Proc. ICLR*, 2020.
- Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin De Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-efficient transfer learning for nlp. In *Proc. ICML*, 2019.
- Wei-Ning Hsu, Benjamin Bolte, Yao-Hung Hubert Tsai, Kushal Lakhotia, Ruslan Salakhutdinov, and Abdelrahman Mohamed. HuBERT: Self-supervised speech representation learning by masked prediction of hidden units. *IEEE/ACM Trans. Audio, Speech, Lang. Process.*, 29:3451–3460, 2021.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *Proc. ICLR*, 2022.
- Ke Hu, Zhehuai Chen, Chao-Han Huck Yang, Piotr Żelasko, Oleksii Hrinchuk, Vitaly Lavrukhin, Jagadeesh Balam, and Boris Ginsburg. Chain-of-thought prompting for speech translation. *arXiv preprint arXiv:2409.11538*, 2024a.
- Shujie Hu, Long Zhou, Shujie Liu, Sanyuan Chen, Lingwei Meng, Hongkun Hao, Jing Pan, Xunying Liu, Jinyu Li, Sunit Sivasankaran, Linqun Liu, and Furu Wei. WavLLM: Towards robust and adaptive speech large language model. In *Findings of EMNLP*, 2024b.
- Chien-Yu Huang, Ke-Han Lu, Shih-Heng Wang, Chi-Yuan Hsiao, Chun-Yi Kuan, Haibin Wu, Siddhant Arora, Kai-Wei Chang, Jiatong Shi, Yifan Peng, Roshan Sharma, Shinji Watanabe, Bhiksha Ramakrishnan, Shady Shehata, and Hung-yi Lee. Dynamic-SUPERB: Towards a dynamic, collaborative, and comprehensive instruction-tuning benchmark for speech. In *Proc. ICASSP*, 2024.
- Chien-Yu Huang et al. Dynamic-SUPERB phase-2: A collaboratively expanding benchmark for measuring the capabilities of spoken language models with 180 tasks. In *Proc. ICLR*, 2025.
- Xuedong Huang, Alex Acero, Hsiao-Wuen Hon, and Raj Reddy. *Spoken language processing: A guide to theory, algorithm, and system development*. Prentice hall PTR, 2001.
- Shengpeng Ji, Yifu Chen, Minghui Fang, Jialong Zuo, Jingyu Lu, Hanting Wang, Ziyue Jiang, Long Zhou, Shujie Liu, Xize Cheng, Xiaoda Yang, Zehan Wang, Qian Yang, Jian Li, Yidi Jiang, Jingzhen He, Yunfei Chu, Jin Xu, and Zhou Zhao. WavChat: A survey of spoken dialogue models. *arXiv preprint arXiv:2411.13577*, 2024.

-
- Carlos E Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. Swe-bench: Can language models resolve real-world github issues? *arXiv preprint arXiv:2310.06770*, 2023.
- Mandar Joshi, Eunsol Choi, Daniel Weld, and Luke Zettlemoyer. TriviaQA: A large scale distantly supervised challenge dataset for reading comprehension. In *Proc. ACL*, 2017.
- Zixuan Ke, Yijia Shao, Haowei Lin, Tatsuya Konishi, Gyuhak Kim, and Bing Liu. Continual pre-training of language models. In *Proc. ICLR*, 2023.
- Emmanuel Keuleers and Marc Brysbaert. Wuggy: A multilingual pseudoword generator. *Behavior research methods*, 42:627–633, 2010.
- Eugene Kharitonov, Jade Copet, Kushal Lakhotia, Tu Anh Nguyen, Paden Tomasello, Ann Lee, Ali Elkahky, Wei-Ning Hsu, Abdelrahman Mohamed, Emmanuel Dupoux, and Yossi Adi. textless-lib: A library for textless spoken language processing. In *Proc. NAACL*, 2022a.
- Eugene Kharitonov, Ann Lee, Adam Polyak, Yossi Adi, Jade Copet, Kushal Lakhotia, Tu-Anh Nguyen, Morgane Rivière, Abdelrahman Mohamed, Emmanuel Dupoux, and Wei-Ning Hsu. Text-free prosody-aware generative spoken language modeling. In *Proc. ACL*, 2022b.
- Eugene Kharitonov, Damien Vincent, Zalán Borsos, Raphaël Marinier, Sertan Girgin, Olivier Pietquin, Matt Sharifi, Marco Tagliasacchi, and Neil Zeghidour. Speak, read and prompt: High-fidelity text-to-speech with minimal supervision. *Trans. ACL*, 11:1703–1718, 2023.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. Overcoming catastrophic forgetting in neural networks. *Proc. NAS*, 114(13):3521–3526, 2017.
- Yuma Koizumi, Kohei Yatabe, Heiga Zen, and Michiel Bacchiani. Wavefit: An iterative and non-autoregressive neural vocoder based on fixed-point iteration. In *Proc. SLT*, 2023.
- Jungil Kong, Jaehyeon Kim, and Jaekyoung Bae. HiFi-GAN: Generative adversarial networks for efficient and high fidelity speech synthesis. In *Proc. NeurIPS*, 2020.
- Felix Kreuk, Adam Polyak, Jade Copet, Eugene Kharitonov, Tu-Anh Nguyen, Morgan Rivière, Wei-Ning Hsu, Abdelrahman Mohamed, Emmanuel Dupoux, and Yossi Adi. Textless speech emotion conversion using discrete & decomposed representations. In *Proc. EMNLP*, 2022.
- Chun-Yi Kuan and Hung-yi Lee. Can large audio-language models truly hear? Tackling hallucinations with multi-task assessment and stepwise audio reasoning. *arXiv preprint arXiv:2410.16130*, 2024.
- Chun-Yi Kuan, Wei-Ping Huang, and Hung-yi Lee. Understanding sounds, missing the questions: The challenge of object hallucination in large audio-language models. *Proc. Interspeech*, 2024.
- Cheng-I Jeff Lai, Yang Zhang, Alexander H. Liu, Shiyu Chang, Yi-Lun Liao, Yung-Sung Chuang, Kaizhi Qian, Sameer Khurana, David D. Cox, and Jim Glass. PARP: Prune, adjust and re-prune for self-supervised speech recognition. In *Proc. NeurIPS*, 2021.
- Cheng-I Jeff Lai, Zhiyun Lu, Liangliang Cao, and Ruoming Pang. Instruction-following speech recognition. *arXiv preprint arXiv:2309.09843*, 2023.
- Kushal Lakhotia, Eugene Kharitonov, Wei-Ning Hsu, Yossi Adi, Adam Polyak, Benjamin Bolte, Tu-Anh Nguyen, Jade Copet, Alexei Baevski, Abdelrahman Mohamed, and Emmanuel Dupoux. On generative spoken language modeling from raw audio. *Trans. ACL*, 9:1336–1354, 2021.
- Siddique Latif, Moazzam Shoukat, Fahad Shamshad, Muhammad Usama, Heriberto Cuayáhuitl, and Björn W. Schuller. Sparks of large audio models: A survey and outlook. *arXiv preprint arXiv:2308.12792*, 2023.

-
- LCM team, Loïc Barrault, Paul-Ambroise Duquenne, Maha Elbayad, Artyom Kozhevnikov, Belen Alastruey, Pierre Andrews, Mariano Coria, Guillaume Couairon, Marta R. Costa-jussà, David Dale, Hady Elsahar, Kevin Heffernan, João Maria Janeiro, Tuan Tran, Christophe Ropers, Eduardo Sánchez, Robin San Roman, Alexandre Mourachko, Safiyyah Saleem, and Holger Schwenk. Large concept models: Language modeling in a sentence representation space. *arXiv preprint arXiv:2412.08821*, 2024.
- Matthew Le, Apoorv Vyas, Bowen Shi, Brian Karrer, Leda Sari, Rashel Moritz, Mary Williamson, Vimal Manohar, Yossi Adi, Jay Mahadeokar, and Wei-Ning Hsu. Voicebox: Text-guided multilingual universal speech generation at scale. In *Proc. NeurIPS*, 2024.
- Ann Lee, Peng-Jen Chen, Changhan Wang, Jiatao Gu, Sravya Popuri, Xutai Ma, Adam Polyak, Yossi Adi, Qing He, Yun Tang, Juan Pino, and Wei-Ning Hsu. Direct speech-to-speech translation with discrete units. In *Proc. ACL*, 2022a.
- Ann Lee, Hongyu Gong, Paul-Ambroise Duquenne, Holger Schwenk, Peng-Jen Chen, Changhan Wang, Sravya Popuri, Yossi Adi, Juan Pino, Jiatao Gu, and Wei-Ning Hsu. Textless speech-to-speech translation on real data. In *Proc. NAACL*, 2022b.
- Doyup Lee, Chiheon Kim, Saehoon Kim, Minsu Cho, and Wook-Shin Han. Autoregressive image generation using residual quantization. In *Proc. CVPR*, 2022c.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven C. H. Hoi. BLIP-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *Proc. ICML*, 2023.
- Guan-Ting Lin, Cheng-Han Chiang, and Hung-yi Lee. Advancing large language models to capture varied speaking styles and respond properly in spoken conversations. In *Proc. ACL*, 2024a.
- Guan-Ting Lin, Prashanth Gurunath Shivakumar, Aditya Gourav, Yile Gu, Ankur Gandhe, Hung-yi Lee, and Ivan Bulyko. Align-SLM: Textless spoken language models with reinforcement learning from AI feedback. *arXiv preprint arXiv:2411.01834*, 2024b.
- Yi-Cheng Lin, Wei-Chih Chen, and Hung-yi Lee. Spoken stereoset: On evaluating social bias toward speaker in speech large language models. In *Proc. SLT*, 2024c.
- Yi-Cheng Lin, Tzu-Quan Lin, Chih-Kai Yang, Ke-Han Lu, Wei-Chih Chen, Chun-Yi Kuan, and Hung-yi Lee. Listen and speak fairly: A study on semantic gender bias in speech integrated large language models. In *Proc. SLT*, 2024d.
- Pengfei Liu, Weizhe Yuan, Jinlan Fu, Zhengbao Jiang, Hiroaki Hayashi, and Graham Neubig. Pre-train, prompt, and predict: A systematic survey of prompting methods in natural language processing. *ACM Comput. Surv.*, 55(9):195:1–195:35, 2023.
- Chen-Chou Lo, Szu-Wei Fu, Wen-Chin Huang, Xin Wang, Junichi Yamagishi, Yu Tsao, and Hsin-Min Wang. MOSNet: Deep learning-based objective assessment for voice conversion. In *Proc. Interspeech*, 2019.
- Ke-Han Lu, Zhehuai Chen, Szu-Wei Fu, He Huang, Boris Ginsburg, Yu-Chiang Frank Wang, and Hung-yi Lee. DeSTA: Enhancing speech language models through descriptive speech-text alignment. In *Proc. Interspeech*, 2024a.
- Ke-Han Lu, Zhehuai Chen, Szu-Wei Fu, Chao-Han Huck Yang, Jagadeesh Balam, Boris Ginsburg, Yu-Chiang Frank Wang, and Hung-yi Lee. Developing instruction-following speech language model without speech instruction-tuning data. *arXiv preprint arXiv:2409.20007*, 2024b.
- Ziyang Ma, Yakun Song, Chenpeng Du, Jian Cong, Zhuo Chen, Yuping Wang, Yuxuan Wang, and Xie Chen. Language model can listen while speaking. *arXiv preprint arXiv:2408.02622*, 2024.
- Gallil Maimon and Yossi Adi. Speaking style conversion in the waveform domain using discrete self-supervised units. In *Findings of EMNLP*, 2023.

-
- Gallil Maimon, Amit Roth, and Yossi Adi. A suite for acoustic language model evaluation. *arXiv preprint arXiv:2409.07437*, 2024.
- Gallil Maimon, Avishai Elmakies, and Yossi Adi. Slamming: Training a speech language model on one gpu in a day. *arXiv preprint arXiv:2502.15814*, 2025.
- Soumi Maiti, Yifan Peng, Shukjae Choi, Jee-weon Jung, Xuankai Chang, and Shinji Watanabe. VoxLM: Unified decoder-only models for consolidating speech recognition, synthesis and speech, text continuation tasks. In *Proc. ICASSP*, 2024.
- Ryo Masumura, Tomohiro Tanaka, Atsushi Ando, Ryo Ishii, Ryuichiro Higashinaka, and Yushi Aono. Neural dialogue context online end-of-turn detection. In *Proc. SIGdial Meeting Disc. Dial.*, 2018.
- Raveesh Meena, Gabriel Skantze, and Joakim Gustafson. Data-driven models for timing feedback responses in a map task dialogue system. *Computer Speech & Language*, 28(4):903–922, 2014.
- Ziqiao Meng, Qichao Wang, Wenqian Cui, Yifei Zhang, Bingzhe Wu, Irwin King, Liang Chen, and Peilin Zhao. Parrot: Autoregressive spoken dialogue language modeling with decoder-only transformers. In *NeurIPS Workshop AI-Driven Speech, Music, and Sound Generation*, 2024.
- Shoval Messica and Yossi Adi. NAST: Noise aware speech tokenization for speech language models. In *Proc. Interspeech*, 2024.
- Microsoft et al. Phi-4-mini technical report: Compact yet powerful multimodal language models via mixture-of-loras. *arXiv preprint arXiv:2503.01743*, 2025.
- Abdelrahman Mohamed, Hung-yi Lee, Lasse Borgholt, Jakob D. Havtorn, Joakim Edin, Christian Igel, Katrin Kirchhoff, Shang-Wen Li, Karen Livescu, Lars Maaloe, Tara N. Sainath, and Shinji Watanabe. Self-supervised speech representation learning: A review. *IEEE J. Sel. Top. Signal Process.*, 16(6):1179–1210, 2022.
- Nasrin Mostafazadeh, Nathanael Chambers, Xiaodong He, Devi Parikh, Dhruv Batra, Lucy Vanderwende, Pushmeet Kohli, and James Allen. A corpus and cloze evaluation for deeper understanding of commonsense stories. In *Proc. NAACL*, 2016.
- Pooneh Mousavi, Luca Della Libera, Jarod Duret, Artem Ploujnikov, Cem Subakan, and Mirco Ravanelli. DASB—discrete audio and speech benchmark. *arXiv preprint arXiv:2406.14294*, 2024a.
- Pooneh Mousavi, Jarod Duret, Salah Zaiem, Luca Della Libera, Artem Ploujnikov, Cem Subakan, and Mirco Ravanelli. How should we extract discrete audio tokens from self-supervised models? In *Proc. Interspeech*, 2024b.
- Eliya Nachmani, Alon Levkovitch, Roy Hirsch, Julian Salazar, Chulayuth Asawaroengchai, Soroosh Mariooryad, Ehud Rivlin, RJ Skerry-Ryan, and Michelle Tadmor Ramanovich. Spoken question answering and speech continuation using spectrogram-powered LLM. In *Proc. ICLR*, 2024.
- Tu Anh Nguyen, Eugene Kharitonov, Jade Copet, Yossi Adi, Wei-Ning Hsu, Ali Elkahky, Paden Tomasello, Robin Algayres, Benoit Sagot, Abdelrahman Mohamed, and Emmanuel Dupoux. Generative spoken dialogue language modeling. *Trans. ACL*, 11:250–266, 2023.
- Tu Anh Nguyen, Benjamin Muller, Bokai Yu, Marta R. Costa-jussa, Maha Elbayad, Sravya Popuri, Christophe Ropers, Paul-Ambroise Duquenne, Robin Algayres, Ruslan Mavlyutov, Itai Gat, Mary Williamson, Gabriel Synnaeve, Juan Pino, Benoît Sagot, and Emmanuel Dupoux. SpiRit-LM: Interleaved spoken and written language model. *Trans. ACL*, 13:30–52, 2025.
- OpenAI. Gpt-4o system card, 2024. URL <https://openai.com/index/gpt-4o-system-card/>.
- OpenAI et al. GPT-4 technical report. *arXiv preprint arXiv:2303.08774*, 2024.

-
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Gray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. In *Proc. NeurIPS*, 2022.
- Jing Pan, Jian Wu, Yashesh Gaur, Sunit Sivasankaran, Zhuo Chen, Shujie Liu, and Jinyu Li. COSMIC: Data efficient instruction-tuning for speech in-context learning. In *Proc. Interspeech*, 2024.
- Se Jin Park, Julian Salazar, Aren Jansen, Keisuke Kinoshita, Yong Man Ro, and RJ Skerry-Ryan. Long-form speech generation with spoken language models. *arXiv preprint arXiv:2412.18603*, 2024.
- Ankita Pasad, Bowen Shi, and Karen Livescu. Comparative layer-wise analysis of self-supervised speech models. In *Proc. ICASSP*, 2023.
- Jing Peng, Yucheng Wang, Yu Xi, Xu Li, Xizhuo Zhang, and Kai Yu. A survey on speech large language models. *arXiv preprint arXiv:2410.18908*, 2024a.
- Yifan Peng, Yui Sudo, Shakeel Muhammad, and Shinji Watanabe. DPHuBERT: Joint distillation and pruning of self-supervised speech models. In *Proc. Interspeech*, 2023a.
- Yifan Peng, Jinchuan Tian, Brian Yan, Dan Berrebbi, Xuankai Chang, Xinjian Li, Jiatong Shi, Siddhant Arora, William Chen, Roshan Sharma, Wangyou Zhang, Yui Sudo, Muhammad Shakeel, Jee weon Jung, Soumi Maiti, and Shinji Watanabe. Reproducing Whisper-style training using an open-source toolkit and publicly available data. In *Proc. ASRU*, 2023b.
- Yifan Peng, Krishna C. Puvvada, Zhehuai Chen, Piotr Zelasko, He Huang, Kunal Dhawan, Ke Hu, Shinji Watanabe, Jagadeesh Balam, and Boris Ginsburg. VoiceTextBlender: Augmenting large language models with speech capabilities via single-stage joint speech-text supervised fine-tuning. *arXiv preprint arXiv:2410.17485*, 2024b.
- Yifan Peng, Yui Sudo, Muhammad Shakeel, and Shinji Watanabe. OWSM-CTC: An open encoder-only speech foundation model for speech recognition, translation, and language identification. In *Proc. ACL*, 2024c.
- Adam Polyak, Yossi Adi, Jade Copet, Eugene Kharonov, Kushal Lakhotia, Wei-Ning Hsu, Abdelrahman Mohamed, and Emmanuel Dupoux. Speech resynthesis from discrete disentangled self-supervised representations. In *Proc. Interspeech*, 2021.
- Krishna C Puvvada, Piotr Żelasko, He Huang, Oleksii Hrinchuk, Nithin Rao Koluguri, Kunal Dhawan, Somshubra Majumdar, Elena Rastorgueva, Zhehuai Chen, Vitaly Lavrukhin, et al. Less is more: Accurate speech recognition & translation without web-scale data. In *Proc. Interspeech*, 2024.
- Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners, 2019. URL https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf.
- Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. Robust speech recognition via large-scale weak supervision. In *Proc. ICML*, 2023.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In *Proc. NeurIPS*, 2023.
- Matthew Roddy, Gabriel Skantze, and Naomi Hartel. Investigating speech features for continuous turn-taking prediction using LSTMs. In *Interspeech*, 2018.
- S Sakshi, Utkarsh Tyagi, Sonal Kumar, Ashish Seth, Ramaneswaran Selvakumar, Oriol Nieto, Ramani Duraiswami, Sreyan Ghosh, and Dinesh Manocha. MMAU: A massive multi-task audio understanding and reasoning benchmark. In *Proc. ICLR*, 2025.

-
- Thomas Schatz, Vijayaditya Peddinti, Francis Bach, Aren Jansen, Hynek Hermansky, and Emmanuel Dupoux. Evaluating speech features with the minimal-pair ABX task: Analysis of the classical MFC/PLP pipeline. In *Proc. Interspeech*, 2013.
- Feiyu Shen, Yiwei Guo, Chenpeng Du, Xie Chen, and Kai Yu. Acoustic BPE for speech generation with discrete tokens. In *Proc. ICASSP*, 2024.
- Jonathan Shen, Ruoming Pang, Ron J Weiss, Mike Schuster, Navdeep Jaitly, Zongheng Yang, Zhifeng Chen, Yu Zhang, Yuxuan Wang, Rj Skerrv-Ryan, et al. Natural tts synthesis by conditioning wavenet on mel spectrogram predictions. In *Proc. ICASSP*, 2018.
- Jiatong Shi, Jinchuan Tian, Yihan Wu, Jee-Weon Jung, Jia Qi Yip, Yoshiki Masuyama, William Chen, Yuning Wu, Yuxun Tang, Massa Baali, Dareen Alharthi, Dong Zhang, Ruifan Deng, Tejes Srivastava, Haibin Wu, Alexander H. Liu, Bhiksha Raj, Qin Jin, Ruihua Song, and Shinji Watanabe. ESPnet-Codec: Comprehensive training and evaluation of neural codecs for audio, music, and speech. In *Proc. SLT*, 2024.
- Suwon Shon, Ankita Pasad, Felix Wu, Pablo Brusco, Yoav Artzi, Karen Livescu, and Kyu J. Han. SLUE: New benchmark tasks for spoken language understanding evaluation on natural speech. In *Proc. ICASSP*, 2022.
- Suwon Shon, Kwangyoung Kim, Yi-Te Hsu, Prashant Sridhar, Shinji Watanabe, and Karen Livescu. DiscreteSLU: A large language model with self-supervised discrete speech units for spoken language understanding. In *Proc. Interspeech*, 2024.
- Yu Shu, Siwei Dong, Guangyao Chen, Wenhao Huang, Ruihua Zhang, Daochen Shi, Qiqi Xiang, and Yemin Shi. LLaSM: Large language and speech model. *arXiv preprint arXiv:2308.15930*, 2023.
- Amitay Sicherman and Yossi Adi. Analysing discrete self supervised speech representation for spoken language modeling. In *Proc. ICASSP*, 2023.
- Hubert Siuzdak, Florian Grötschla, and Luca A Lanzendörfer. Snac: Multi-scale neural audio codec. *arXiv preprint arXiv:2410.14411*, 2024.
- Gabriel Skantze. Towards a general, continuous model of turn-taking in spoken dialogue using LSTM recurrent neural networks. In *Proc. SIGdial Meeting Disc. and Dial.*, 2017.
- Aarohi Srivastava et al. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *Trans. MLR*, 2023.
- Changli Tang, Wenyi Yu, Guangzhi Sun, Xianzhao Chen, Tian Tan, Wei Li, Lu Lu, Zejun MA, and Chao Zhang. SALMONN: Towards generic hearing abilities for large language models. In *Proc. ICLR*, 2024.
- Jinchuan Tian, Jiatong Shi, William Chen, Siddhant Arora, Yoshiki Masuyama, Takashi Maekaku, Yihan Wu, Junyi Peng, Shikhar Bharadwaj, Yiwen Zhao, et al. Espnet-speechlm: An open speech language model toolkit. *arXiv preprint arXiv:2502.15218*, 2025.
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, Aurelien Rodriguez, Armand Joulin, Edouard Grave, and Guillaume Lample. LLaMA: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023a.
- Hugo Touvron et al. LLaMA 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023b.
- Emiru Tsunoo, Hayato Futami, Yosuke Kashiwagi, Siddhant Arora, and Shinji Watanabe. Decoder-only architecture for streaming end-to-end speech recognition. In *Proc. Interspeech*, 2024.
- Arnon Turetzky and Yossi Adi. LAST: Language model aware speech tokenization. *arXiv preprint arXiv:2409.03701*, 2024.

-
- Aaron van den Oord, Oriol Vinyals, and Koray Kavukcuoglu. Neural discrete representation learning. In *Proc. NeurIPS*, 2017.
- Bandhav Veluri, Benjamin N Peloquin, Bokai Yu, Hongyu Gong, and Shyamnath Gollakota. Beyond turn-based interfaces: Synchronous LLMs as full-duplex dialogue agents. In *Proc. EMNLP*, 2024.
- Bin Wang, Xunlong Zou, Geyu Lin, Shuo Sun, Zhuohan Liu, Wenyu Zhang, Zhengyuan Liu, AiTi Aw, and Nancy F. Chen. AudioBench: A universal benchmark for audio large language models. *arXiv preprint arXiv:2406.16020*, 2024a.
- Chen Wang, Minpeng Liao, Zhongqiang Huang, Jinliang Lu, Junhong Wu, Yuchen Liu, Chengqing Zong, and Jiajun Zhang. BLSP: Bootstrapping language-speech pre-training via behavior alignment of continuation writing. *arXiv preprint arXiv:2309.00916*, 2023a.
- Mingqiu Wang, Wei Han, Izhak Shafran, Zelin Wu, Chung-Cheng Chiu, Yuan Cao, Yongqiang Wang, Nanxin Chen, Yu Zhang, Hagen Soltau, Paul Rubenstein, Lukas Zilka, Dian Yu, Zhong Meng, Golan Pundak, Nikhil Siddhartha, Johan Schalkwyk, and Yonghui Wu. SLM: Bridge the thin gap between speech and text foundation models. In *Proc. ASRU*, 2023b.
- Peng Wang, Songshuo Lu, Yaohua Tang, Sijie Yan, Wei Xia, and Yuanjun Xiong. A full-duplex speech dialogue scheme based on large language models. In *Proc. NeurIPS*, 2024b.
- Xiong Wang, Yangze Li, Chaoyou Fu, Yunhang Shen, Lei Xie, Ke Li, Xing Sun, and Long Ma. Freeze-Omni: A smart and low latency speech-to-speech dialogue model with frozen LLM. *arXiv preprint arXiv:2411.00774*, 2024c.
- Alex Warstadt, Alicia Parrish, Haokun Liu, Anhad Mohananey, Wei Peng, Sheng-Fu Wang, and Samuel R. Bowman. BLiMP: The benchmark of linguistic minimal pairs for English. *Trans. ACL*, 8:377–392, 2020.
- Jason Wei, Maarten Bosma, Vincent Y. Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M. Dai, and Quoc V. Le. Finetuned language models are zero-shot learners. In *Proc. ICLR*, 2022a.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *Proc. NeurIPS*, 2022b.
- WhisperSpeech. WhisperSpeech: An open source text-to-speech system built by inverting Whisper, 2024. URL <https://github.com/collabora/WhisperSpeech>.
- Felix Wu, Kwangyoun Kim, Shinji Watanabe, Kyu J. Han, Ryan McDonald, Kilian Q. Weinberger, and Yoav Artzi. Wav2Seq: Pre-training speech-to-text encoder-decoder models using pseudo languages. In *Proc. ICASSP*, 2023a.
- Haibin Wu, Xuanjun Chen, Yi-Cheng Lin, Kai-wei Chang, Ho-Lam Chung, Alexander H. Liu, and Hung-yi Lee. Towards audio language modeling-an overview. *arXiv preprint arXiv:2402.13236*, 2024a.
- Haibin Wu, Xuanjun Chen, Yi-Cheng Lin, Kaiwei Chang, Jiawei Du, Ke-Han Lu, Alexander H. Liu, Ho-Lam Chung, Yuan-Kuei Wu, Dongchao Yang, Songxiang Liu, Yi-Chiao Wu, Xu Tan, James Glass, Shinji Watanabe, and Hung-yi Lee. Codec-Superb @ SLT 2024: A lightweight benchmark for neural audio codec models. In *Proc. SLT*, 2024b.
- Haibin Wu, Ho-Lam Chung, Yi-Cheng Lin, Yuan-Kuei Wu, Xuanjun Chen, Yu-Chi Pai, Hsiu-Hsuan Wang, Kai-Wei Chang, Alexander H. Liu, and Hung-yi Lee. Codec-SUPERB: An in-depth analysis of sound codec models. In *Findings of ACL*, 2024c.
- Haibin Wu, Naoyuki Kanda, Sefik Emre Eskimez, and Jinyu Li. TS3-Codec: Transformer-based simple streaming single codec. *arXiv preprint arXiv:2411.18803*, 2024d.
- Haibin Wu, Yuan Tseng, and Hung-yi Lee. CodecFake: Enhancing anti-spoofing models against deepfake audios from codec-based speech synthesis systems. In *Proc. Interspeech*, 2024e.

-
- Jian Wu, Yashesh Gaur, Zhuo Chen, Long Zhou, Yimeng Zhu, Tianrui Wang, Jinyu Li, Shujie Liu, Bo Ren, Linqun Liu, and Yu Wu. On decoder-only architecture for speech-to-text and large language model integration. In *Proc. ASRU*, 2023b.
- Zhifei Xie and Changqiao Wu. Mini-Omni: Language models can hear, talk while thinking in streaming. *arXiv preprint arXiv:2408.16725*, 2024a.
- Zhifei Xie and Changqiao Wu. Mini-Omni2: Towards open-source GPT-4o model with vision, speech and duplex. *arXiv preprint arXiv:2410.11190*, 2024b.
- Detai Xin, Xu Tan, Shinnosuke Takamichi, and Hiroshi Saruwatari. BigCodec: Pushing the limits of low-bitrate neural speech codec. *arXiv preprint arXiv:2409.05377*, 2024.
- Wang Xu, Shuo Wang, Weilin Zhao, Xu Han, Yukun Yan, Yudi Zhang, Zhe Tao, Zhiyuan Liu, and Wanxiang Che. Enabling real-time conversations with minimal training costs. *arXiv preprint arXiv:2409.11727*, 2024.
- Hongfei Xue, Yuhao Liang, Bingshen Mu, Shiliang Zhang, Mengzhe Chen, Qian Chen, and Lei Xie. E-chat: Emotion-sensitive spoken dialogue system with large language models. In *Proc. ISCSLP*, 2024.
- Dongchao Yang, Jinchuan Tian, Xu Tan, Rongjie Huang, Songxiang Liu, Haohan Guo, Xuankai Chang, Jiatong Shi, Sheng Zhao, Jiang Bian, Zhou Zhao, Xixin Wu, and Helen M. Meng. UniAudio: Towards universal audio generation with large language models. In *Proc. ICML*, 2024a.
- Qian Yang, Jin Xu, Wenrui Liu, Yunfei Chu, Ziyue Jiang, Xiaohuan Zhou, Yichong Leng, Yuanjun Lv, Zhou Zhao, Chang Zhou, and Jingren Zhou. AIR-Bench: Benchmarking large audio-language models via generative comprehension. In *Proc. ACL*, 2024b.
- Lili Yu, Daniel Simig, Colin Flaherty, Armen Aghajanyan, Luke Zettlemoyer, and Mike Lewis. MEGABYTE: Predicting million-byte sequences with multiscale transformers. In *Proc. NeurIPS*, 2023.
- Wenyi Yu, Changli Tang, Guangzhi Sun, Xianzhao Chen, Tian Tan, Wei Li, Lu Lu, Zejun Ma, and Chao Zhang. Connecting speech encoder and large language model for ASR. In *Proc. ICASSP*, 2024.
- Ze Yuan, Yanqing Liu, Shujie Liu, and Sheng Zhao. Continuous speech tokens makes llms robust multi-modality learners. *arXiv preprint arXiv:2412.04917*, 2024.
- Neil Zeghidour, Alejandro Luebs, Ahmed Omran, Jan Skoglund, and Marco Tagliasacchi. SoundStream: An end-to-end neural audio codec. *IEEE Trans. Audio, Speech, Lang. Process.*, 30:495–507, 2022.
- Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. HellaSwag: Can a machine really finish your sentence? In *Proc. ACL*, 2019.
- Aohan Zeng, Zhengxiao Du, Mingdao Liu, Kedong Wang, Shengmin Jiang, Lei Zhao, Yuxiao Dong, and Jie Tang. GLM-4-Voice: Towards intelligent and human-like end-to-end spoken chatbot. *arXiv preprint arXiv:2412.02612*, 2024a.
- Aohan Zeng, Zhengxiao Du, Mingdao Liu, Lei Zhang, Shengmin Jiang, Yuxiao Dong, and Jie Tang. Scaling speech-text pre-training with synthetic interleaved data. *arXiv preprint arXiv:2411.17607*, 2024b.
- Jun Zhan, Junqi Dai, Jiasheng Ye, Yunhua Zhou, Dong Zhang, Zhigeng Liu, Xin Zhang, Ruibin Yuan, Ge Zhang, Linyang Li, Hang Yan, Jie Fu, Tao Gui, Tianxiang Sun, Yu-Gang Jiang, and Xipeng Qiu. AnyGPT: Unified multimodal LLM with discrete sequence modeling. In *Proc. ACL*, 2024.
- Dong Zhang, Shimin Li, Xin Zhang, Jun Zhan, Pengyu Wang, Yaqian Zhou, and Xipeng Qiu. SpeechGPT: Empowering large language models with intrinsic cross-modal conversational abilities. In *Findings of EMNLP*, 2023a.
- Duzhen Zhang, Yahan Yu, Jiahua Dong, Chenxing Li, Dan Su, Chenhui Chu, and Dong Yu. MM-LLMs: Recent advances in multimodal large language models. In *Findings of ACL*, 2024a.

-
- Qinglin Zhang, Luyao Cheng, Chong Deng, Qian Chen, Wen Wang, Siqi Zheng, Jiaqing Liu, Hai Yu, Chaohong Tan, Zhihao Du, and Shiliang Zhang. OmniFlatten: An end-to-end GPT model for seamless voice conversation. *arXiv preprint arXiv:2410.17799*, 2025.
- Susan Zhang, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuohui Chen, Christopher Dewan, Mona Diab, Xian Li, Xi Victoria Lin, Todor Mihaylov, Myle Ott, Sam Shleifer, Kurt Shuster, Daniel Simig, Punit Singh Koura, Anjali Sridhar, Tianlu Wang, and Luke Zettlemoyer. OPT: Open pre-trained transformer language models. *arXiv preprint arXiv:2205.01068*, 2022.
- Xin Zhang, Dong Zhang, Shimin Li, Yaqian Zhou, and Xipeng Qiu. SpeechTokenizer: Unified speech tokenizer for speech language models. In *Proc. ICLR*, 2024b.
- Xinrong Zhang, Yingfa Chen, Shengding Hu, Xu Han, Zihang Xu, Yuanwei Xu, Weilin Zhao, Maosong Sun, and Zhiyuan Liu. Beyond the turn-based game: Enabling real-time conversations with duplex models. In *Proc. EMNLP*, 2024c.
- Yu Zhang, Wei Han, James Qin, Yongqiang Wang, Ankur Bapna, Zhehuai Chen, Nanxin Chen, Bo Li, Vera Axelrod, Gary Wang, Zhong Meng, Ke Hu, Andrew Rosenberg, Rohit Prabhavalkar, Daniel S. Park, Parisa Haghani, Jason Riesa, Ginger Perng, Hagen Soltau, Trevor Strohman, Bhuvana Ramabhadran, Tara N. Sainath, Pedro J. Moreno, Chung-Cheng Chiu, Johan Schalkwyk, Françoise Beaufays, and Yonghui Wu. Google USM: Scaling automatic speech recognition beyond 100 languages. *arXiv preprint arXiv:2303.01037*, 2023b.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-judge with MT-bench and Chatbot Arena. In *NeurIPS Datasets and Benchmarks Track*, 2023.

A Appendix

Table 3: Spoken language models. *P.*: Phonetic token. *A.*: Acoustic token. *C.*: Continuous feature. For components not explicitly cited in the table, please see the following references: SSE Algayres et al. (2022), USM (Zhang et al., 2023b), SNAC (Siuzdak et al., 2024), SenseVoice (An et al., 2024), and Canary (Puvvada et al., 2024) for speech representation models, and Tacotron 2 (Shen et al., 2018), WaveFit (Koizumi et al., 2023), and CosyVoice (Du et al., 2024) for speech decoders. Modalities other than text or speech (e.g., vision) are omitted. [†]: Pre-trained ASR encoder.

Model	Input Representation → Output Representation	Speech Decoder	Modality Adapter
<i>Pure Speech LM</i>			
GSLM (Lakhotia et al., 2021)	HuBERT (<i>P.</i>)	Tacotron 2	X
pGSLM (Kharitonov et al., 2022b)	HuBERT (<i>P.</i>) + F0 + duration	HiFi-GAN	X
dGSLM (Nguyen et al., 2023)	HuBERT (<i>P.</i>) + duration	HiFi-GAN	X
tGSLM (Algayres et al., 2023)	SSE (<i>P.</i>) + HuBERT (<i>P.</i>)	Tacotron 2	X
AudioLM (Borsos et al., 2023a)	w2v-BERT (<i>P.</i>) → SoundStream (<i>A.</i>)	SoundStream	X
TWIST (Hassid et al., 2023)	HuBERT (<i>P.</i>)	HiFi-GAN	X
SoundStorm (Borsos et al., 2023b)	w2v-BERT (<i>P.</i>) → SoundStream (<i>A.</i>)	SoundStream	X
Align-SLM (Lin et al., 2024b)	HuBERT(<i>P.</i>)	HiFi-GAN	X
<i>Speech+Text LM</i>			
AnyGPT (Zhan et al., 2024)	Text + SpeechTokenizer (<i>A.</i>)	SpeechTokenizer	Vocabulary Expansion
SpeechGPT (Zhang et al., 2023a)	Text + mHuBERT (<i>P.</i>)	HiFi-GAN	Vocabulary Expansion
VoxLM (Maiti et al., 2024)	Text + HuBERT (<i>P.</i>)	HiFi-GAN	Vocabulary Expansion
Spectron (Nachmani et al., 2024)	Text + Mel-Spectrogram (<i>C.</i>)	WaveFit	MLP
SpiRit-LM (Nguyen et al., 2025)	Text + HuBERT (<i>P.</i>)	HiFi-GAN	Vocabulary Expansion
Moshi (Défossez et al., 2024)	Text + Mimi (<i>A.</i>)	Mimi	Vocabulary Expansion
MiniOmni (Xie & Wu, 2024a)	Text + Whisper [†] (<i>C.</i>) → Text + SNAC (<i>A.</i>)	SNAC	MLP & Vocabulary Expansion
LLaMA-Omni (Fang et al., 2025)	Whisper (<i>C.</i>) → Text + HuBERT (<i>P.</i>)	HiFi-GAN	MLP
SLAM-Omni (Chen et al., 2024b)	Text + Whisper [†] (<i>P.</i>)	CosyVoice	Linear
GLM-4-Voice (Zeng et al., 2024a)	Text + Whisper [†] (<i>P.</i>)	CosyVoice	Vocabulary Expansion
MinMo (Chen et al., 2025a)	Text + SenseVoice [†] (<i>C.</i>) → Text + SenseVoice [†] (<i>P.</i>)	CosyVoice	Transformer + CNN
<i>Speech-Aware Text LM</i>			
WavPrompt (Gao et al., 2022)	Text + Wav2vec 2.0 (<i>C.</i>) → Text	X	CNN
X-LLM (Chen et al., 2023)	Text + Conformer [†] (<i>C.</i>) → Text	X	CIF + Transformer
LTU-AS (Gong et al., 2023b)	Text + Whisper [†] (<i>C.</i>) → Text	X	TLTR
COSMIC (Pan et al., 2024)	Text + Whisper [†] (<i>C.</i>) → Text	X	Window-level Q-Former
SALMONN (Tang et al., 2024)	Text + Whisper [†] (<i>C.</i>) + BEATS (<i>C.</i>) → Text	X	Window-level Q-Former
Speech-LLaMA (Wu et al., 2023b)	Mel-Spectrogram(<i>C.</i>) → Text	X	CTC compressor + Transformer
SLM (Wang et al., 2023b)	Text + USM [†] (<i>C.</i>) → Text	X	Transformer
LLaSM (Shu et al., 2023)	Text + Whisper [†] (<i>C.</i>) → Text	X	Linear
SALM (Chen et al., 2024d)	Text + Conformer Encoder [†] (<i>C.</i>) → Text	X	Conformer
Qwen-Audio (Chu et al., 2023)	Text + Whisper [†] (<i>C.</i>) → Text	X	Pooling layer
DiscreteSLU (Shon et al., 2024)	Text + WavLM (<i>P.</i>) → Text	X	Embedding + CNN + Transformer + Linear
SpeechVerse (Das et al., 2024)	Text + Whisper [†] (<i>C.</i>) → Text	X	CNN
DeSTA (Lu et al., 2024a)	Text + Whisper [†] (<i>C.</i>) → Text	X	Q-Former
WavLLM (Hu et al., 2024b)	Text + Whisper [†] (<i>C.</i>) + WavLM (<i>C.</i>) → Text	X	CNN + Bottleneck adapter + Linear
VoiceTextBlender (Peng et al., 2024b)	Text + Canary Encoder [†] (<i>C.</i>) → Text	X	Conformer
Phi-4-Multimodal (Microsoft et al., 2025)	Text + Conformer Encoder(<i>C.</i>) → Text	X	MLP




Figure 6: Development timeline of spoken language models. “Publicly available” refers to models with publicly released weights (but not necessarily code, data, or other artifacts). “Commercial SLM” include some systems that handle additional (non-linguistic) modalities, such as images. For more details and references, see the main article and Table 3.

Table 4: A representative set of training and evaluation strategies for speech-aware text language models, along with key tasks, details on the training data, and key insights for further context.

Model	Training Strategy	Training Tasks	Evaluation Tasks	Training Data	Findings
SLM Wang et al. (2023b)	IT	ASR, ST, Alpaca tasks	ASR, ST, contextual biasing, open-ended QA	multilingual ASR, ST datasets, Alpaca with TTS (93k hrs)	Fine-tuning only a lightweight adapter is sufficient.
SALMONN Tang et al. (2024)	Pre-training → IT	ASR, AAC, 15 audio and speech tasks	8 seen audio and speech task + 7 unseen ST and SLU tasks	ASR, ST and SLU datasets (4k hrs)	Performs complex reasoning tasks like audio-based storytelling in zero-shot setting.
LTU-AS Gong et al. (2023b)	3-stage training	Classification, general QA	audio and speech tasks, open-ended QA	Open-ASQA dataset (9.6M QA pairs)	Single IF model on both speech and audio tasks is feasible.
COSMIC Pan et al. (2024)	IT	ASR, QA	ASR, SQA, ST, contextual biasing, ASR (out of domain)	TED-LIUM 3 (452 hours)	Shows few-shot in-context learning ability.
LTU Gong et al. (2024)	4-stage training	Classification + cesc., close-ended QA, general QA	Classification, captioning, open-ended QA	AQA (5M QA pairs)	Performs <i>open-ended</i> audio tasks.
SALM Chen et al. (2024d)	IT	ASR, ST	ASR, ST, keyword boosting	LibriSpeech (960h), IWSLT 2023 (4.8k hrs)	Biases the model to predict keywords in instruction.
Qwen-Audio Chu et al. (2023)	Pre-training → IT	speech, audio and music tasks, audio dialogue	Speech, audio and music tasks, qualitative examples of audio analysis/editing	No details (30k hrs for ASR, >123k hrs in total)	Performs chat-based training to learn conversational ability.
DeSTA2 Lu et al. (2024b)	IT	Audio captioning	Dynamic-SUPERB Huang et al. (2024), AIR-Bench Yang et al. (2024b)	Mixture of several datasets (155 hours)	Training only on audio captioning can generalize to other tasks
DiscreteSLU (Shon et al., 2024)	IT	ASR, SQA, sentiment analysis, NER	WER, S2ST	Tedlium-3 (Hernandez et al., 2018) & SLUE (Shon et al., 2022)	Discrete speech token input can be competitive with continuous representations, in both seen and unseen speech domains, even in a zero-shot setting.